Cairo University

Faculty of Engineering

Electronics and Electrical Communications Engineering

Assignment 3

ELC 4028 – Artificial Neural Networks

ID	**Section**	**Name**
9220399	2	Shahd Gamal Mahmoud Hussein
9220411	2	Salah El-Din Hassan Hassan Mohamed
9220984	4	Yousef Khaled Omar Mahmoud

Supervised by:

Dr. Mohsen Rashwan
Dr. Mohamed Motaz

Contents

1 Problem 1 VAE Synthetic Data with Low-Data Stabilization 3
1.1 Conditional Variational Autoencoder (VAE) . . . 3
1.1.1 Pipeline Overview . . . 3
1.1.2 Pipeline Flowchart . . . 5
1.1.3 Original MNIST Samples . . . 6
1.1.4 Data Augmentation . . . 6
1.1.5 VAE Architecture . . . 6
1.1.6 Generated Samples During Training . . . 7
1.1.7 Synthetic Dataset Construction . . . 8
1.1.8 Set A Samples . . . 8
1.1.9 Set B Samples . . . 8
1.1.10 Set C Samples . . . 9
1.1.11 Experimental Results . . . 9
1.1.12 Discussion . . . 9
1.1.13 Conclusion . . . 10

2 Problem 2 GAN Synthetic Data with Low-Data Stabilization 11
2.1 Conditional Generative Adversarial Network (GAN) . . . 11
2.1.1 Pipeline Overview . . . 11
2.1.2 Pipeline Flowchart . . . 13
2.1.3 Augmented Training Samples . . . 14
2.1.4 GAN Architecture . . . 14
2.1.5 GAN Training Curves . . . 14
2.1.6 Discriminator Mistake Analysis . . . 17
2.1.7 Generated GAN Samples . . . 19
2.1.8 Synthetic Dataset Construction . . . 19
2.1.9 Set A Samples . . . 20
2.1.10 Set B Samples . . . 20
2.1.11 Set C Samples . . . 20
2.1.12 Experimental Results . . . 20
2.1.13 Summary Table for Problem 1 and Problem 2 . . . 21
2.1.14 Discussion . . . 21
2.1.15 Conclusion . . . 22

3 Problem 3 Impact of Attention Mechanisms 23
3.1 Part (a): CNN vs. Attention CNN on ReducedMNIST . . . 23
3.1.1 Introduction . . . 23
3.1.2 Dataset . . . 23
3.1.3 Network Architectures . . . 23
3.1.4 Training Configuration . . . 24
3.1.5 Results . . . 25
3.1.6 Analysis . . . 27
3.1.7 Insights and Observations . . . 27
3.1.8 Suggestions for Future Improvements . . . 28
3.2 Part (b): CNN vs. Attention CNN on Spoken Digits (FSDD) . . . 29
3.2.1 Introduction . . . 29

3.2.2 Dataset and Preprocessing . . . 29
3.2.3 Network Architectures . . . 30
3.2.4 Training Configuration . . . 31
3.2.5 Results . . . 33
3.2.6 Analysis . . . 34
3.2.7 Insights and Observations . . . 35
3.2.8 Suggestions for Future Improvements . . . 35
3.3 Cross-Part Comparison: When Does Attention Help? . . . 36

4 Problem 4 General Information Retrieval 37
4.1 System Overview . . . 37
4.2 Selected Book . . . 37
4.3 Preprocessing and Chunking . . . 37
4.4 Embedding and Indexing . . . 38
4.5 Language Model . . . 38
4.6 Retrieval Methods . . . 38
4.7 System Implementation . . . 38
4.7.1 `main.py` — Core Pipeline . . . 38
4.7.2 `web_ui.py` — Interactive Web Interface . . . 39
4.7.3 Web UI Screenshots . . . 40
4.8 Results from 10 Arabic Queries (Classical vs. Semantic) . . . 41
4.9 Answers for 10 Queries: RAG vs. LLM-only . . . 46
4.10 Comparison and Reflection . . . 48
4.10.1 Classical vs. Semantic Retrieval . . . 48
4.10.2 RAG vs. LLM-only Generation . . . 49
4.10.3 Conclusions . . . 49
4.11 Tools and Libraries Used . . . 50

5 Problem 5 Benchmarking Arabic NLP Tasks 51
5.1 Task Definition and Scope . . . 51
5.2 Dataset Creation and Cleaning . . . 51
5.3 Annotation Process . . . 52
5.4 Evaluation Setup . . . 53
5.4.1 Jais first-attempt failure and prompt correction . . . 53
5.5 Quantitative Results . . . 54
5.6 Selected Output Examples . . . 55
5.7 Qualitative Observations . . . 56
5.8 Error Analysis . . . 56
5.8.1 Arabic-specific challenge 1: dialect variation . . . 57
5.8.2 Arabic-specific challenge 2: morphology and orthographic variation 57
5.8.3 Arabic-specific challenge 3: ambiguity and context dependence . . . 57
5.8.4 Arabic-specific challenge 4: Classical Arabic satire . . . 57
5.9 Difficult and Borderline Cases . . . 58
5.10 Limitations . . . 58
5.11 Final Conclusions and Recommendations . . . 58
5.12 Checklist . . . 59
5.13 Submitted Files for Problem 5 . . . 60

List of Figures

1.1 Colored pipeline of the Conditional VAE synthetic data generation framework. . . . 5
1.2 Original MNIST samples. . . . 6
1.3 Examples of augmented MNIST samples using rotation, shifting, scaling, and noise injection. . . . 6
1.4 Generated samples from VAE Run 1. . . . 7
1.5 Generated samples from VAE Run 2. . . . 7
1.6 Generated samples from VAE Run 3. . . . 7
1.7 Generated samples from VAE Run 4. . . . 7
1.8 Generated samples from VAE Run 5. . . . 8
1.9 Samples from Set A (all generated samples). . . . 8
1.10 Samples from Set B (high-confidence samples). . . . 8
1.11 Samples from Set C (mid-confidence samples). . . . 9
2.1 Conditional GAN synthetic data generation pipeline. . . . 13
2.2 Augmented MNIST samples used for GAN training. . . . 14
2.3 GAN loss curves for Run 1. . . . 15
2.4 GAN loss curves for Run 2. . . . 15
2.5 GAN loss curves for Run 3. . . . 16
2.6 GAN loss curves for Run 4. . . . 16
2.7 GAN loss curves for Run 5. . . . 17
2.8 Discriminator mistakes for GAN Run 1. . . . 17
2.9 Discriminator mistakes for GAN Run 2. . . . 18
2.10 Discriminator mistakes for GAN Run 3. . . . 18
2.11 Discriminator mistakes for GAN Run 4. . . . 18
2.12 Discriminator mistakes for GAN Run 5. . . . 19
2.13 Generated samples from GAN Run 1. . . . 19
2.14 Samples from Set A (all generated samples). . . . 20
2.15 Samples from Set B (high-confidence samples). . . . 20
2.16 Samples from Set C (mid-confidence samples). . . . 20
2.17 Classifier accuracy comparison using GAN-generated datasets. . . . 21
3.1 Baseline CNN training log (ReducedMNIST). . . . 25
3.2 Attention CNN training log (ReducedMNIST). . . . 25
3.3 Baseline CNN accuracy and loss curves (ReducedMNIST). . . . 25
3.4 Attention CNN accuracy and loss curves (ReducedMNIST). . . . 26
3.5 Terminal output of the comprehensive model comparison (ReducedMNIST). 26
3.6 Baseline CNN training log (FSDD). . . . 32
3.7 Attention CNN training log (FSDD). . . . 32
3.8 Baseline CNN accuracy and loss curves (FSDD). . . . 33
3.9 Attention CNN accuracy and loss curves (FSDD). . . . 33
3.10 Terminal output of the comprehensive model comparison (FSDD). . . . 34
4.1 Landing Page — Main Web UI Interface 40
4.2 Classical vs. Semantic Retrieval Example — Side-by-Side Comparison . . . 40
4.3 RAG vs. LLM-only Example — Grounded vs. Hallucinated Answers 41
5.1 F1 score comparison across the five evaluated LLMs. . . . 54
5.2 F1 score by Arabic variety group. . . . 55

1 Problem 1 VAE Synthetic Data with Low-Data Stabilization

1.1 Conditional Variational Autoencoder (VAE)

In this experiment, a Conditional Variational Autoencoder (CVAE) was implemented to generate synthetic MNIST handwritten digit images using only a limited real dataset of 350 samples per digit.

The main objective was to investigate whether synthetic data generated by a VAE can improve the performance of a LeNet-5 classifier trained on a small dataset.

1.1.1 Pipeline Overview

The implemented pipeline is summarized as follows:

1. Load and preprocess MNIST dataset.
2. Reduce the training dataset to only 350 real samples per digit.
3. Apply data augmentation using:
   - Rotation
   - Translation
   - Scaling
   - Gaussian noise
4. Train a Conditional VAE using the augmented dataset.
5. Repeat VAE training for 5 independent runs using different random weight initialization.
6. Generate 1000 synthetic images per digit from each VAE run.
7. Train a LeNet-5 classifier using only the 350 real samples.
8. Classify generated samples and compute:

$$\text{confidence} = \max(\text{softmax output})$$

9. Create three synthetic datasets:
   - Set A: all generated samples
   - Set B: confidence ≥ 0.9
   - Set C: $0.6 \leq$ confidence < 0.9
10. Retrain LeNet-5 using:
    - 350 real samples only
    - 350 real + Set A
    - 350 real + Set B
    - 350 real + Set C

11. Compare against:
    - 350-real baseline
    - 1000-real baseline

1.1.2 Pipeline Flowchart

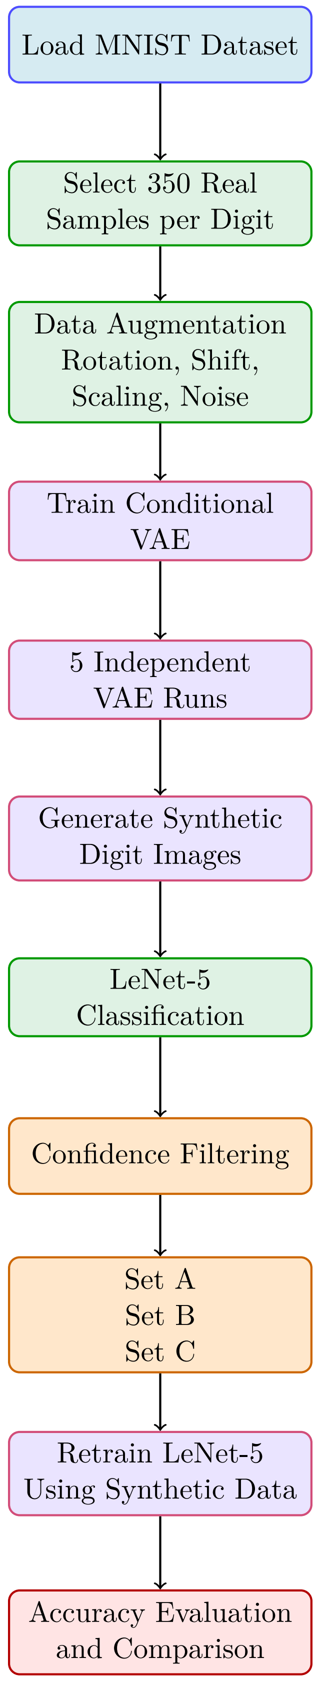

Figure 1.1: Colored pipeline of the Conditional VAE synthetic data generation framework.

1.1.3 Original MNIST Samples

Figure 1.2 shows examples from the original MNIST dataset.

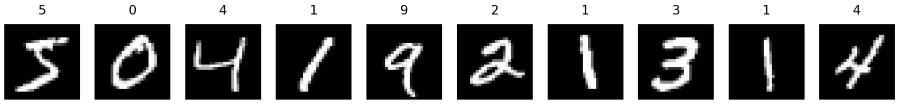

Figure 1.2: Original MNIST samples.

1.1.4 Data Augmentation

To compensate for the limited dataset size, multiple augmentation techniques were applied.

These augmentations improve generalization and help the VAE learn digit variations. Figure 1.3 shows augmented examples.

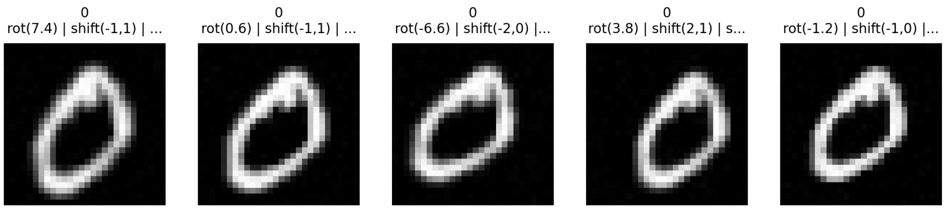

Figure 1.3: Examples of augmented MNIST samples using rotation, shifting, scaling, and noise injection.

1.1.5 VAE Architecture

The implemented Conditional VAE consists of:

- Encoder:
  - Convolutional layers
  - Latent mean and variance estimation
- Decoder:
  - Dense projection
  - Reshape layer
  - Transposed convolution layers
  - Sigmoid output layer

The latent dimension was selected as:

$$z \in \mathbb{R}^8$$

The loss function combines:

- Binary Cross Entropy reconstruction loss
- KL-divergence regularization

1.1.6 Generated Samples During Training

Five independent VAE runs were trained using different random initialization.
Figures 1.4 – 1.8 show generated samples from different runs.

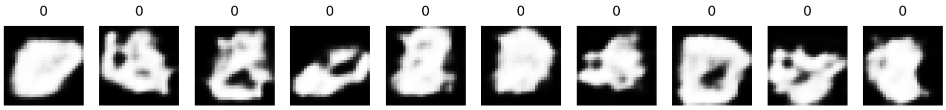

Figure 1.4: Generated samples from VAE Run 1.

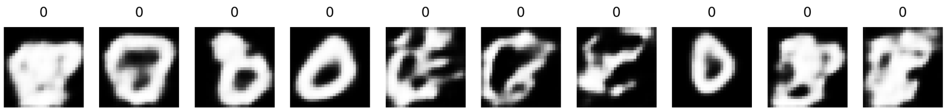

Figure 1.5: Generated samples from VAE Run 2.

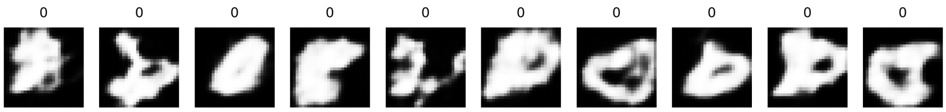

Figure 1.6: Generated samples from VAE Run 3.

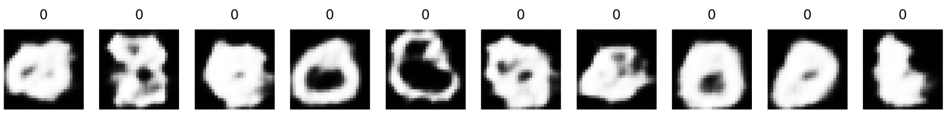

Figure 1.7: Generated samples from VAE Run 4.

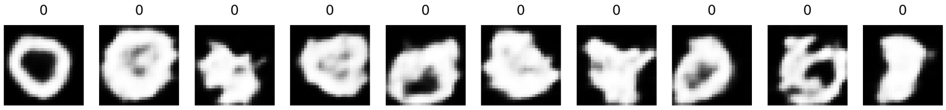

Figure 1.8: Generated samples from VAE Run 5.

The generated images demonstrate that the Conditional VAE successfully learned the structural characteristics of handwritten digits while maintaining variation across samples.

1.1.7 Synthetic Dataset Construction

Generated samples were filtered using the trained LeNet-5 classifier confidence.

Three datasets were created:

- **Set A**: all generated samples
- **Set B**: high-confidence samples
- **Set C**: mid-confidence samples

1.1.8 Set A Samples

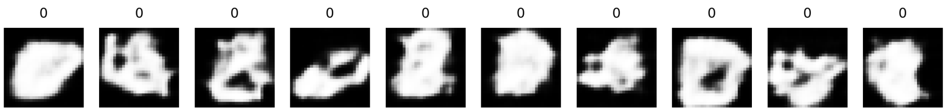

Figure 1.9: Samples from Set A (all generated samples).

Set A contains all generated images without filtering. Some samples are realistic while others contain distortions and noisy structures.

1.1.9 Set B Samples

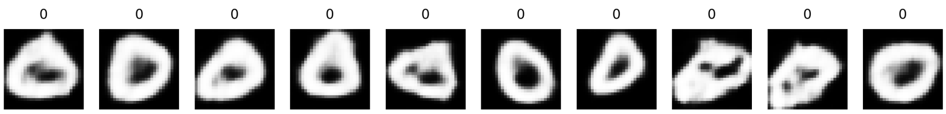

Figure 1.10: Samples from Set B (high-confidence samples).

Set B contains cleaner and more realistic samples because only high-confidence predictions were accepted.

1.1.10 Set C Samples

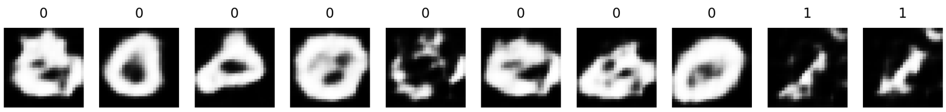

Figure 1.11: Samples from Set C (mid-confidence samples).

Set C contains moderately confident samples that preserve diversity while removing highly corrupted outputs.

1.1.11 Experimental Results

The final LeNet-5 classification accuracies are shown in Table 1.1.

Table 1.1: Comparison of classifier accuracy using different synthetic datasets.

Experiment	**Accuracy**
350 Real Baseline	96.36%
Set A	93.31%
Set B	95.25%
Set C	95.29%
1000 Real Baseline	98.17%

1.1.12 Discussion

The experimental results show that:

- Using only 350 real samples achieved a strong baseline accuracy of 96.36%.
- Adding all generated samples (Set A) reduced performance due to the inclusion of low-quality and distorted generated images.
- Confidence filtering significantly improved performance.
- Set B and Set C achieved accuracies close to the real-only baseline.
- High-confidence filtering successfully removed unrealistic samples while preserving useful digit variations.
- The 1000-real baseline still achieved the best performance because real samples preserve the true MNIST distribution more accurately than synthetic samples.

1.1.13 Conclusion

A Conditional Variational Autoencoder was successfully implemented for MNIST handwritten digit generation.

The generated samples demonstrated that the VAE learned meaningful latent representations despite training on a very limited dataset.

Experimental evaluation showed that synthetic data quality strongly affects classifier performance. Confidence-based filtering improved the usefulness of generated samples and produced classification accuracy close to the real-only baseline.

The results indicate that VAE-generated data can partially compensate for limited datasets, although synthetic samples still do not fully replace real training data.

2 Problem 2 GAN Synthetic Data with Low-Data Stabilization

2.1 Conditional Generative Adversarial Network (GAN)

In this experiment, a Conditional Generative Adversarial Network (CGAN) was implemented to generate synthetic MNIST handwritten digit images using a limited dataset containing only 350 real samples per digit.

The main objective was to investigate whether GAN-generated synthetic data can improve classifier performance and reduce the dependence on large real datasets.

In addition, the experiment compares GAN-generated data against traditional augmentation techniques used in Assignment 2.

2.1.1 Pipeline Overview

The implemented GAN pipeline is summarized as follows:

1. Load and preprocess the MNIST dataset.
2. Reduce the dataset to 350 real samples per digit.
3. Apply augmentation using:
   - Rotation
   - Translation
   - Scaling
   - Gaussian noise
4. Train a Conditional GAN using the augmented dataset.
5. Repeat GAN training for 5 independent runs using different random initialization.
6. Generate synthetic digit samples from each GAN model.
7. Train a LeNet-5 classifier using the reduced real dataset.
8. Classify generated images and compute:

$$\text{confidence} = \max(\text{softmax output})$$

9. Construct three synthetic datasets:
   - Set A: all generated samples
   - Set B: confidence ≥ 0.9
   - Set C: $0.6 \leq$ confidence < 0.9
10. Retrain LeNet-5 using:
    - 350 real samples only
    - 350 real + Set A
    - 350 real + Set B

- 350 real + Set C

11. Compare against:

    - 350-real baseline
    - 1000-real baseline
    - Augmentation-only results from Assignment 2

2.1.2 Pipeline Flowchart

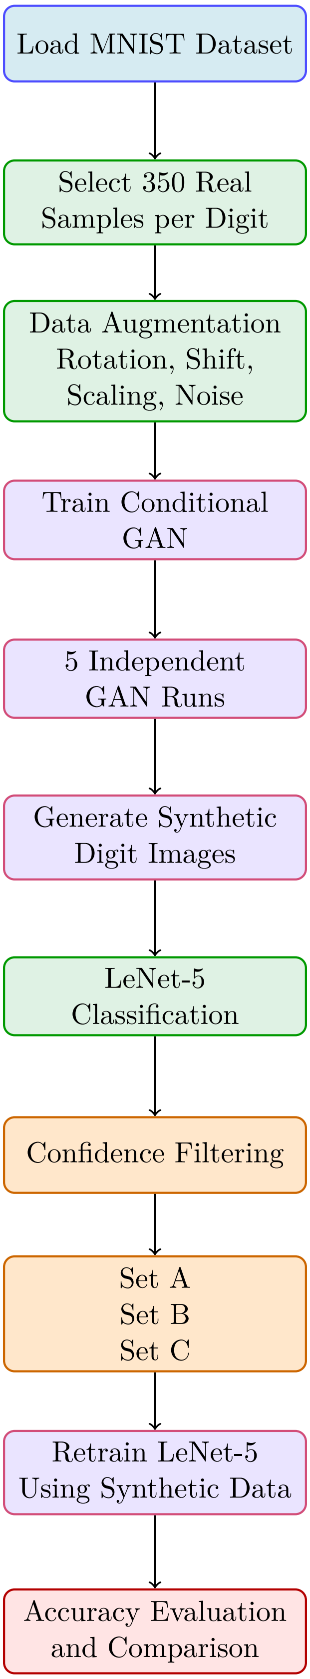

Figure 2.1: Conditional GAN synthetic data generation pipeline.

2.1.3 Augmented Training Samples

Figure 2.2 shows examples of augmented MNIST samples used during GAN training.

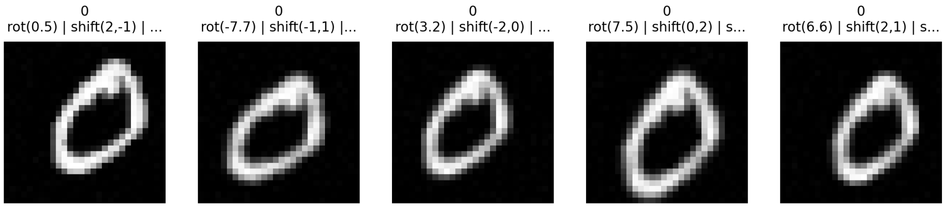

Figure 2.2: Augmented MNIST samples used for GAN training.

2.1.4 GAN Architecture

The implemented Conditional GAN consists of:

- **Generator**
  - Latent noise input
  - Label embedding
  - Dense projection
  - Transposed convolution layers
  - Tanh output activation
- **Discriminator**
  - Convolutional feature extraction
  - Label conditioning
  - LeakyReLU activations
  - Binary classification output

The GAN was trained using adversarial learning where:

- The generator attempts to fool the discriminator.
- The discriminator attempts to distinguish real from fake images.

2.1.5 GAN Training Curves

Figures 2.3 – 2.7 show discriminator and generator losses for the five independent GAN runs.

Figure 2.3: GAN loss curves for Run 1.

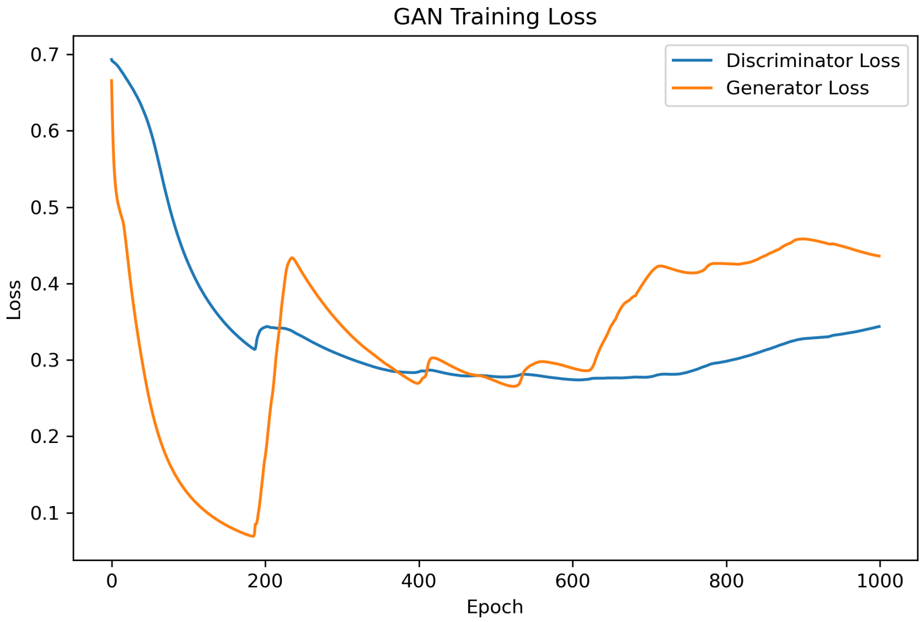

Figure 2.4: GAN loss curves for Run 2.

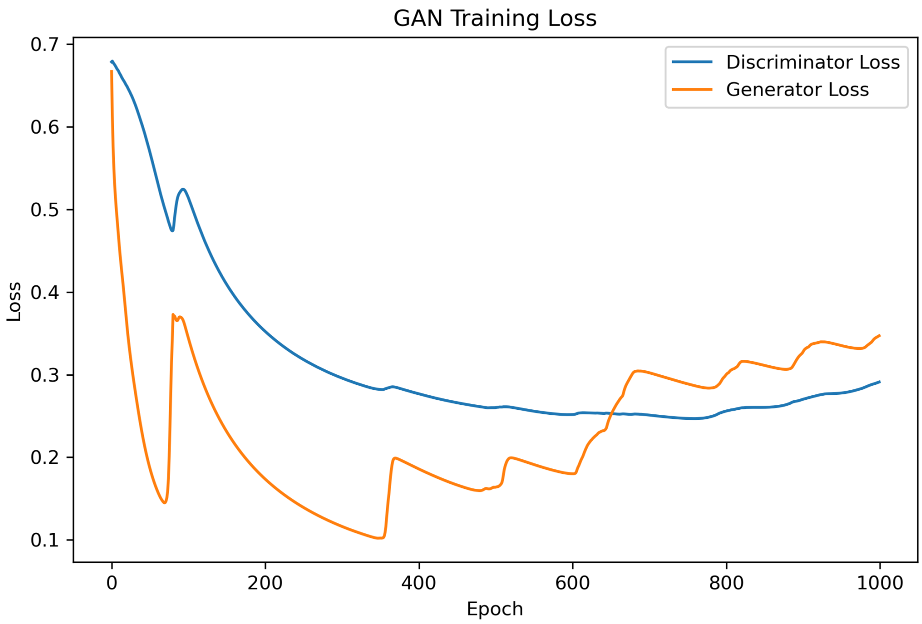

Figure 2.5: GAN loss curves for Run 3.

Figure 2.6: GAN loss curves for Run 4.

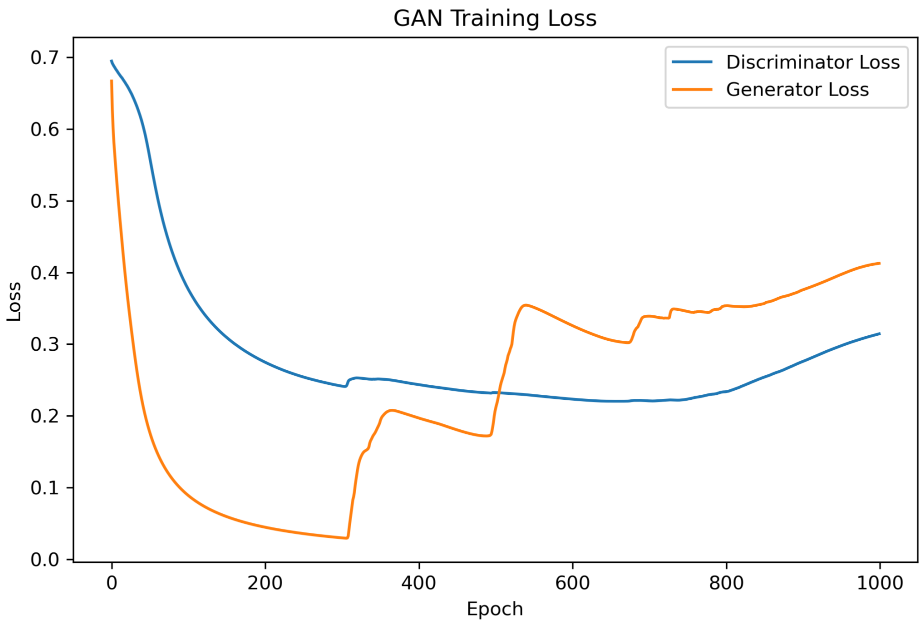

Figure 2.7: GAN loss curves for Run 5.

The loss curves demonstrate unstable adversarial behavior during training. Generator and discriminator losses oscillate significantly across runs, which indicates unstable convergence.

2.1.6 Discriminator Mistake Analysis

Figures 2.8 – 2.12 show discriminator classification mistakes.

Green labels indicate fake images classified as real.

Red labels indicate real images classified as fake.

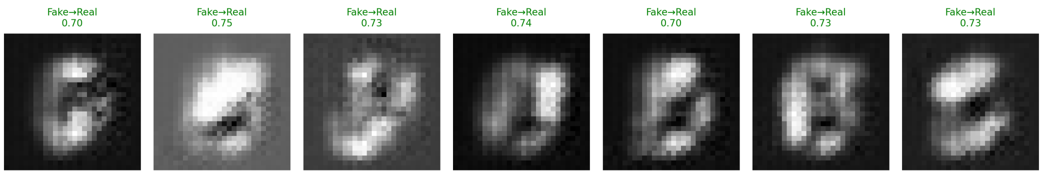

Figure 2.8: Discriminator mistakes for GAN Run 1.

Figure 2.9: Discriminator mistakes for GAN Run 2.

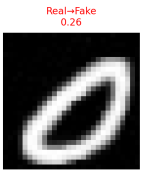

Figure 2.10: Discriminator mistakes for GAN Run 3.

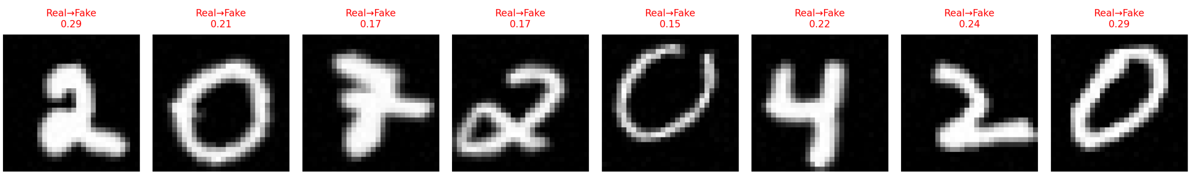

Figure 2.11: Discriminator mistakes for GAN Run 4.

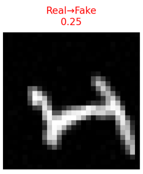

Figure 2.12: Discriminator mistakes for GAN Run 5.

The discriminator mistakes show that some generated samples successfully fooled the discriminator despite containing visible distortions and unrealistic structures.

2.1.7 Generated GAN Samples

Figure 2.13 shows generated samples from one GAN run.

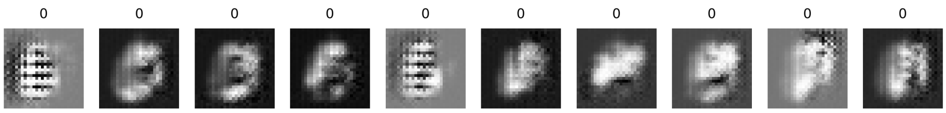

Figure 2.13: Generated samples from GAN Run 1.

The generated digits demonstrate that the GAN learned general digit structure. However, many samples remain blurry or contain artifacts.

Some generated samples appear realistic, while others collapse into noisy or distorted patterns.

2.1.8 Synthetic Dataset Construction

Generated samples were filtered using LeNet-5 classifier confidence.

Three datasets were constructed:

- **Set A**: all generated samples
- **Set B**: confidence ≥ 0.9
- **Set C**: $0.6 \leq$ confidence < 0.9

2.1.9 Set A Samples

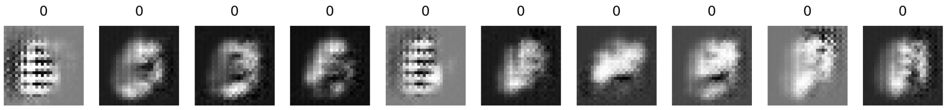

Figure 2.14: Samples from Set A (all generated samples).

Set A includes all generated images. Many samples contain severe artifacts and unstable structures.

2.1.10 Set B Samples

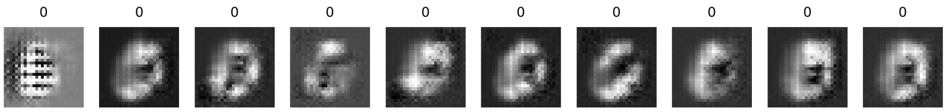

Figure 2.15: Samples from Set B (high-confidence samples).

Set B contains higher-quality samples selected using confidence filtering.

2.1.11 Set C Samples

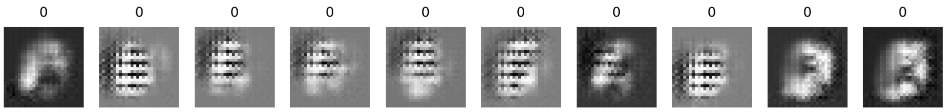

Figure 2.16: Samples from Set C (mid-confidence samples).

Set C preserves moderately confident samples while removing highly corrupted outputs.

2.1.12 Experimental Results

Figure 2.17 summarizes classifier performance using GAN-generated synthetic datasets.

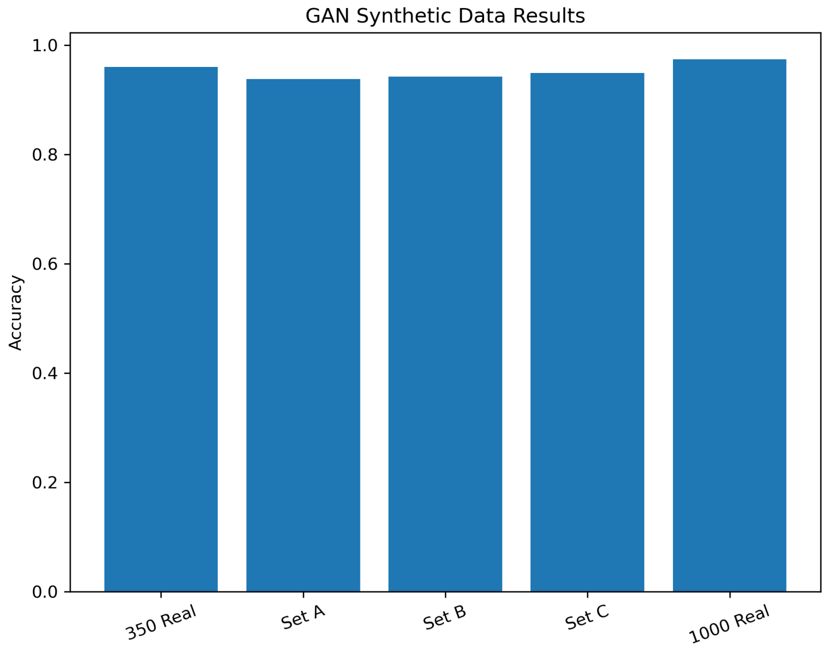

Figure 2.17: Classifier accuracy comparison using GAN-generated datasets.

2.1.13 Summary Table for Problem 1 and Problem 2

Table 2.1 summarizes the results of both VAE and GAN approaches.

Table 2.1: Comparison of VAE and GAN synthetic data strategies.

Model Used	No. of Examples	Confidence Level	Accuracy
VAE	All generated data	Set A	93.31%
VAE	CL > 0.9	Set B	95.25%
VAE	0.6 < CL < 0.9	Set C	95.29%
GAN	All generated data	Set A	93.8%
GAN	CL > 0.9	Set B	94.4%
GAN	0.6 < CL < 0.9	Set C	94.8%
350 Real Baseline	Real only	—	96.36%
1000 Real Baseline	Real only	—	98.17%

2.1.14 Discussion

The experimental results demonstrate several important observations:

- GAN training was significantly more unstable than VAE training.
- The adversarial nature of GAN optimization caused oscillations between generator and discriminator performance.
- Generated GAN samples varied greatly in quality across runs.
- Some generated digits appeared realistic, while others were blurry or highly distorted.

- Confidence filtering improved the usefulness of generated samples by removing low-quality outputs.
- Set B and Set C consistently performed better than Set A because filtering removed unrealistic samples.
- Despite improvements, GAN-generated data still did not outperform the real-only baseline.
- Compared to Assignment 2 augmentation, traditional augmentation remained more stable and reliable.
- The experiments show that synthetic data can partially compensate for limited datasets, but fully replacing real data remains difficult.

2.1.15 Conclusion

A Conditional GAN was successfully implemented for MNIST handwritten digit generation using a highly limited dataset.

The generated samples demonstrated that the GAN learned meaningful digit structures; however, training instability caused large quality variations across runs.

Visual inspection confirmed that some generated samples were realistic while others suffered from mode collapse, blur, and severe artifacts.

Experimental evaluation showed that confidence-based filtering improved classifier performance by selecting more reliable synthetic samples.

Overall, GAN-generated data partially reduced the need for real data, but the results indicate that synthetic GAN samples still cannot fully replace real training samples.

Compared to VAE generation and traditional augmentation, GAN training was less stable and more sensitive to initialization and adversarial dynamics.

3 Problem 3 Impact of Attention Mechanisms

3.1 Part (a): CNN vs. Attention CNN on ReducedMNIST

3.1.1 Introduction

This part builds two CNN classifiers on the ReducedMNIST dataset — a subset of the standard MNIST handwritten-digit benchmark — and measures the impact of inserting a spatial attention mechanism. The two models are:

1. **Baseline CNN** — a compact two-block convolutional network with no attention, similar in spirit to LeNet-5.
2. **Attention CNN** — the same architecture augmented with a `SpatialAttention` layer after each convolutional block.

3.1.2 Dataset

ReducedMNIST is a down-sampled version of MNIST consisting of 28×28 greyscale images of handwritten digits (0–9). The dataset is split into training and validation sets, with inputs normalised to $[0, 1]$.

3.1.3 Network Architectures

Baseline CNN The baseline model is a compact sequential network. Table 3.1 shows its layer structure extracted from the model summary.

Table 3.1: Baseline CNN architecture for ReducedMNIST

Layer (type)	Output Shape	Param #
Conv2D 3×3, 4 filters, ReLU	(None, 26, 26, 4)	40
MaxPooling2D 2×2	(None, 13, 13, 4)	0
Conv2D 3×3, 8 filters, ReLU	(None, 11, 11, 8)	296
MaxPooling2D 2×2	(None, 5, 5, 8)	0
Flatten	(None, 200)	0
Dense (60, ReLU)	(None, 60)	12,060
Dense (10, Softmax)	(None, 10)	610
Total trainable parameters		**13,006**
Non-trainable parameters		0

Spatial Attention Mechanism The `SpatialAttention` layer generates a 2D attention mask from the feature map by applying average-pooling and max-pooling along the channel dimension, concatenating the results, and passing them through a 3×3 (or 7×7) sigmoid convolution:

$$\mathbf{M} = \sigma(\mathrm{Conv}([\,\mathrm{AvgPool}_C(\mathbf{F});\ \mathrm{MaxPool}_C(\mathbf{F})\,]))\,, \qquad \mathbf{F}' = \mathbf{F} \odot \mathbf{M} \tag{1}$$

Each `SpatialAttention` module adds **19 parameters**: $(3 \times 3 \times 2) + 1 = 19$.

Attention CNN The attention model inserts a `SpatialAttention` layer immediately after each convolutional layer, before pooling. Table 3.2 shows the full structure.

Table 3.2: Attention CNN architecture for ReducedMNIST

Layer (type)	Output Shape	Param #
Conv2D 3×3, 4 filters, ReLU	(None, 26, 26, 4)	40
SpatialAttention	(None, 26, 26, 4)	**19**
MaxPooling2D 2×2	(None, 13, 13, 4)	0
Conv2D 3×3, 8 filters, ReLU	(None, 11, 11, 8)	296
SpatialAttention	(None, 11, 11, 8)	**19**
MaxPooling2D 2×2	(None, 5, 5, 8)	0
Flatten	(None, 200)	0
Dense (60, ReLU)	(None, 60)	12,060
Dense (10, Softmax)	(None, 10)	610
Total trainable parameters		**13,044**
Non-trainable parameters		0

Parameter Overhead of Attention Two `SpatialAttention` modules add $2 \times 19 = +\mathbf{38}$ parameters on top of 13,006 — an overhead of $\approx 0.29\%$.

3.1.4 Training Configuration

Both models are trained under identical conditions summarised in Table 3.3.

Table 3.3: Training hyperparameters for Part (a)

Hyperparameter	Value
Optimizer	Adam
Loss function	Sparse categorical cross-entropy
Epochs	10
Input image size	28 × 28 × 1
Number of classes	10 (digits 0–9)

```
=== Training Baseline Model ===
Epoch 1/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 8ms/step - accuracy: 0.7072 - loss: 1.0206 - val_accuracy: 0.9067 - val_loss: 0.3319
Epoch 2/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9177 - loss: 0.2743 - val_accuracy: 0.9351 - val_loss: 0.2184
Epoch 3/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9415 - loss: 0.1979 - val_accuracy: 0.9471 - val_loss: 0.1759
Epoch 4/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9509 - loss: 0.1617 - val_accuracy: 0.9524 - val_loss: 0.1551
Epoch 5/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9574 - loss: 0.1401 - val_accuracy: 0.9608 - val_loss: 0.1307
Epoch 6/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9631 - loss: 0.1221 - val_accuracy: 0.9622 - val_loss: 0.1227
Epoch 7/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9651 - loss: 0.1100 - val_accuracy: 0.9633 - val_loss: 0.1145
Epoch 8/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9697 - loss: 0.0987 - val_accuracy: 0.9667 - val_loss: 0.1103
Epoch 9/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9726 - loss: 0.0890 - val_accuracy: 0.9599 - val_loss: 0.1247
Epoch 10/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9746 - loss: 0.0832 - val_accuracy: 0.9665 - val_loss: 0.1032
```

Figure 3.1: Baseline CNN training log (ReducedMNIST).

```
=== Training Attention Model ===
Epoch 1/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 3s 13ms/step - accuracy: 0.6444 - loss: 1.3104 - val_accuracy: 0.8793 - val_loss: 0.4231
Epoch 2/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 11ms/step - accuracy: 0.9062 - loss: 0.3283 - val_accuracy: 0.9288 - val_loss: 0.2490
Epoch 3/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 13ms/step - accuracy: 0.9348 - loss: 0.2216 - val_accuracy: 0.9419 - val_loss: 0.1977
Epoch 4/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 14ms/step - accuracy: 0.9487 - loss: 0.1719 - val_accuracy: 0.9528 - val_loss: 0.1589
Epoch 5/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 12ms/step - accuracy: 0.9576 - loss: 0.1467 - val_accuracy: 0.9552 - val_loss: 0.1474
Epoch 6/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 12ms/step - accuracy: 0.9632 - loss: 0.1251 - val_accuracy: 0.9608 - val_loss: 0.1259
Epoch 7/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 11ms/step - accuracy: 0.9653 - loss: 0.1144 - val_accuracy: 0.9621 - val_loss: 0.1317
Epoch 8/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 11ms/step - accuracy: 0.9703 - loss: 0.1022 - val_accuracy: 0.9471 - val_loss: 0.1791
Epoch 9/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 14ms/step - accuracy: 0.9743 - loss: 0.0881 - val_accuracy: 0.9667 - val_loss: 0.1143
Epoch 10/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 3s 16ms/step - accuracy: 0.9752 - loss: 0.0796 - val_accuracy: 0.9658 - val_loss: 0.1166
```

Figure 3.2: Attention CNN training log (ReducedMNIST).

Training Logs

3.1.5 Results

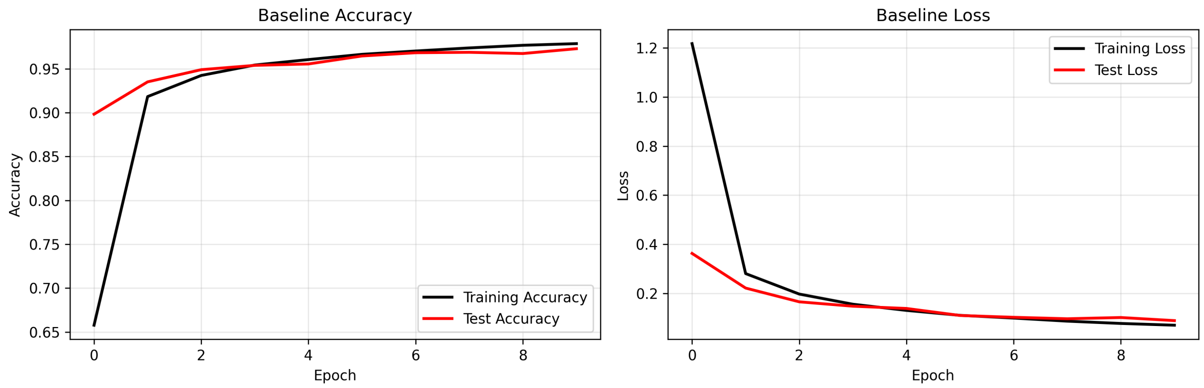

Figure 3.3: Baseline CNN accuracy and loss curves (ReducedMNIST).

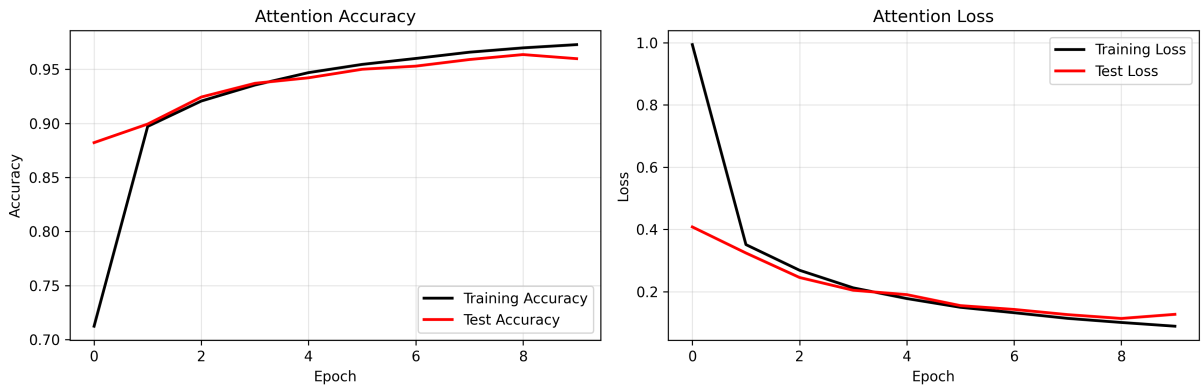

Figure 3.4: Attention CNN accuracy and loss curves (ReducedMNIST).

Training and Validation Curves

Table 3.4: Final-epoch performance: Baseline vs. Attention CNN (ReducedMNIST)

Metric	Baseline CNN	Attention CNN	Difference
Train Accuracy (%)	97.46	97.52	+0.06
Test Accuracy (%)	96.65	96.58	−0.07
Train Loss	0.0832	0.0796	−0.0036
Test Loss	0.1032	0.1166	+0.0134
Training Time (s)	11.75	21.65	+9.90
Train-Test Acc. Gap	0.81%	0.94%	+0.13 pp
Model Parameters	13,006	13,044	+38

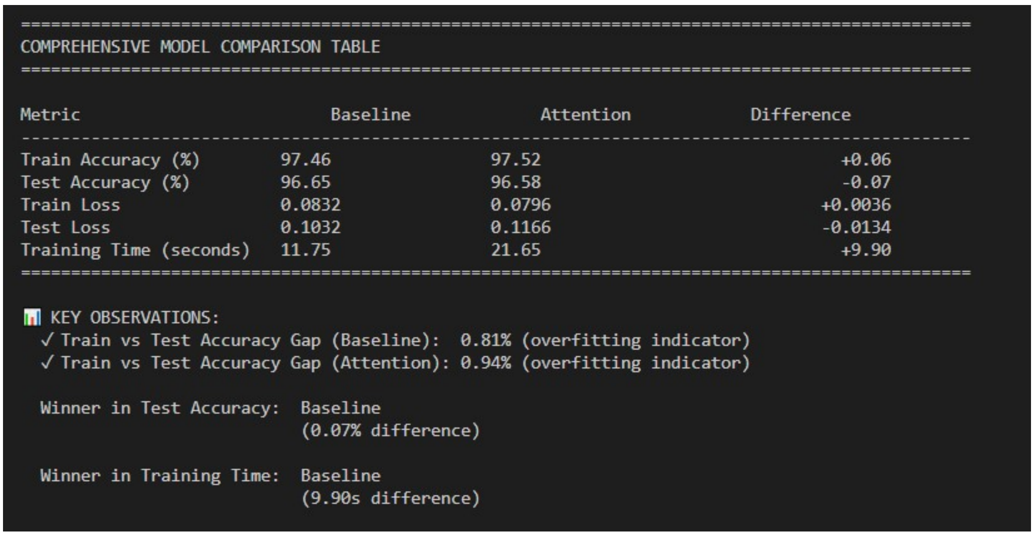

```
==========================================================================================
COMPREHENSIVE MODEL COMPARISON TABLE
==========================================================================================

Metric                         Baseline            Attention           Difference
------------------------------------------------------------------------------------------
Train Accuracy (%)        97.46               97.52                              +0.06
Test Accuracy (%)         96.65               96.58                              -0.07
Train Loss                0.0832              0.0796                           +0.0036
Test Loss                 0.1032              0.1166                           -0.0134
Training Time (seconds)   11.75               21.65                              +9.90
==========================================================================================

📊 KEY OBSERVATIONS:
  ✓ Train vs Test Accuracy Gap (Baseline):  0.81% (overfitting indicator)
  ✓ Train vs Test Accuracy Gap (Attention): 0.94% (overfitting indicator)

  Winner in Test Accuracy:  Baseline
                            (0.07% difference)

  Winner in Training Time:  Baseline
                            (9.90s difference)
```

Figure 3.5: Terminal output of the comprehensive model comparison (ReducedMNIST).

Quantitative Comparison

Key Observations

- **Near-identical accuracy:** Both models achieve ≈96.6% test accuracy. The difference (−0.07 pp) is negligible and likely within run-to-run variance.
- **Winner in test accuracy: Baseline** (96.65% vs. 96.58%).
- **Winner in training time: Baseline** (11.75 s vs. 21.65 s) — the attention module roughly doubles per-epoch cost on this small model.
- **Generalisation gap:** Both models generalise well (gap < 1%), indicating that ReducedMNIST is straightforward enough that attention provides no regularisation benefit.
- **Minimal parameter overhead:** Only +38 weights (≈ 0.29%), yet the time cost is disproportionately high due to the extra conv operations on every forward pass.

3.1.6 Analysis

Why Attention Makes Little Difference Here MNIST digits are centred, well-segmented, and relatively noise-free. The background pixels carry almost no useful signal, and a simple CNN already learns to ignore them through weight sharing and pooling. Spatial attention, which is designed to *suppress* uninformative regions, therefore finds little to improve upon — the baseline already focuses on the right pixels implicitly.

Convergence Behaviour Both models converge smoothly within the first 2–3 epochs (Figures 3.3 and 3.4), with training and validation curves tracking each other closely throughout all 10 epochs. There is no sign of overfitting, which confirms that the dataset is not challenging enough to reveal a difference between the two architectures.

Training Time Overhead The attention model takes ≈21.65 s vs. 11.75 s for the baseline — nearly double. This is because the `SpatialAttention` module performs two pooling operations and one additional convolution on every batch, which is relatively costly when the base model is already very lightweight (only 13,006 parameters).

3.1.7 Insights and Observations

1. **Attention is not universally beneficial.** On clean, centred datasets like ReducedMNIST, the baseline already extracts sufficient features, and attention adds complexity without improving accuracy.
2. **Dataset difficulty matters.** Attention mechanisms tend to shine on harder tasks where informative regions are small, sparse, or mixed with strong background noise.
3. **Parameter efficiency.** The 38-parameter overhead is tiny, but the time cost is not — this illustrates that parameter count alone is not a reliable proxy for computational cost.
4. **Both models generalise well.** Train-test gaps of < 1% indicate no meaningful overfitting on ReducedMNIST.

3.1.8 Suggestions for Future Improvements

- **Test on harder datasets** (e.g. Fashion-MNIST or CIFAR-10) where attention mechanisms are more likely to provide measurable gains.
- **Channel attention (SE blocks):** On MNIST, selecting *which feature maps* to emphasise may be more useful than selecting spatial regions.
- **More training epochs:** Both accuracy curves are still slightly rising at epoch 10; extending to 20–30 epochs with early stopping may reveal a difference.
- **Lighter attention modules:** A 1×1 conv instead of 3×3 would reduce the time overhead while preserving the attention mechanism.

3.2 Part (b): CNN vs. Attention CNN on Spoken Digits (FSDD)

3.2.1 Introduction

This part revisits the spoken-digit classification task from Assignment 2. Two CNN architectures are trained on mel-spectrogram images of spoken digits from the Free Spoken Digit Dataset (FSDD):

1. **Baseline CNN** — a two-block convolutional classifier with no attention mechanism.
2. **Attention CNN** — the same architecture augmented with a `SpatialAttentionWithMask` layer after each convolutional block.

3.2.2 Dataset and Preprocessing

Free Spoken Digit Dataset (FSDD) The FSDD contains 3,000 audio recordings (300 per digit, 10 digits), recorded at 8 kHz by multiple speakers.

Spectrogram Generation

1. Audio loaded at $f_s = 8{,}000\,\text{Hz}$ using `librosa`.
2. Mel-spectrogram: $N_{\text{FFT}} = 512$, hop = 256, $n_{\text{mels}} = 64$.
3. Converted to dB: $S_{\text{dB}} = 10 \log_{10}(S/S_{\max})$.
4. Min-max normalised and resized to 64×64 pixels.
5. Channel dimension appended: final shape (64, 64, 1).

Train/Test Split A stratified 80/20 split (`random_state = 42`) gives approximately 2,400 training and 600 test samples, balanced across all ten digit classes.

3.2.3 Network Architectures

Table 3.5: Baseline CNN architecture for FSDD

Layer (type)	Output Shape	Param #
Conv2D 3×3, 32 filters, ReLU	(None, 64, 64, 32)	320
BatchNormalization	(None, 64, 64, 32)	128
MaxPooling2D 2×2	(None, 32, 32, 32)	0
Dropout (0.25)	(None, 32, 32, 32)	0
Conv2D 3×3, 64 filters, ReLU	(None, 32, 32, 64)	18,496
BatchNormalization	(None, 32, 32, 64)	256
MaxPooling2D 2×2	(None, 16, 16, 64)	0
Dropout (0.30)	(None, 16, 16, 64)	0
Flatten	(None, 16,384)	0
Dense (256, ReLU)	(None, 256)	4,194,560
Dense (10, Softmax)	(None, 10)	2,570
Total trainable parameters		**4,216,330**

Baseline CNN

Spatial Attention Mechanism The `SpatialAttentionWithMask` layer computes:

$$\mathbf{M} = \sigma(\mathrm{Conv}_{7\times7}([\,\mathrm{AvgPool}_C(\mathbf{F});\ \mathrm{MaxPool}_C(\mathbf{F})\,]))\,, \qquad \mathbf{F}' = \mathbf{F} \odot \mathbf{M} \tag{2}$$

Each module uses a 7×7 Conv2D with 2 input channels and 1 output channel: $(7 \times 7 \times 2) + 1 = \mathbf{99}$ parameters.

Table 3.6: Attention CNN architecture for FSDD

Layer (type)	Output Shape	Param #
InputLayer	(None, 64, 64, 1)	0
Conv2D 3×3, 32 filters, ReLU	(None, 64, 64, 32)	320
BatchNormalization	(None, 64, 64, 32)	128
SpatialAttentionWithMask (attn_1)	[(None,64,64,32), (None,64,64,1)]	**99**
MaxPooling2D 2×2	(None, 32, 32, 32)	0
Conv2D 3×3, 64 filters, ReLU	(None, 32, 32, 64)	18,496
BatchNormalization	(None, 32, 32, 64)	256
SpatialAttentionWithMask (attn_2)	[(None,32,32,64), (None,32,32,1)]	**99**
MaxPooling2D 2×2	(None, 16, 16, 64)	0
Flatten	(None, 16,384)	0
Dense (256, ReLU)	(None, 256)	4,194,560
Dense (10, Softmax)	(None, 10)	2,570
Total trainable parameters		**4,216,528**

Attention CNN

Parameter Overhead of Attention Two attention modules add $2 \times 99 = +198$ parameters over the baseline, an increase of only $\approx 0.005\%$.

3.2.4 Training Configuration

Table 3.7: Training hyperparameters for Part (b)

Hyperparameter	Value
Optimizer	Adam (lr = 0.001)
Loss function	Sparse categorical cross-entropy
Batch size	32
Epochs	10
Train / Test split	80% / 20% (stratified)
Input image size	$64 \times 64 \times 1$
Number of classes	10 (digits 0–9)
Sample rate	8,000 Hz
FFT window (N_{FFT})	512
Hop length	256
Mel bins	64

```
=== Training Baseline Model ===
Epoch 1/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 8ms/step - accuracy: 0.7072 - loss: 1.0206 - val_accuracy: 0.9067 - val_loss: 0.3319
Epoch 2/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9177 - loss: 0.2743 - val_accuracy: 0.9351 - val_loss: 0.2184
Epoch 3/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9415 - loss: 0.1979 - val_accuracy: 0.9471 - val_loss: 0.1759
Epoch 4/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9509 - loss: 0.1617 - val_accuracy: 0.9524 - val_loss: 0.1551
Epoch 5/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9574 - loss: 0.1401 - val_accuracy: 0.9608 - val_loss: 0.1307
Epoch 6/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9631 - loss: 0.1221 - val_accuracy: 0.9622 - val_loss: 0.1227
Epoch 7/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9651 - loss: 0.1100 - val_accuracy: 0.9633 - val_loss: 0.1145
Epoch 8/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 6ms/step - accuracy: 0.9697 - loss: 0.0987 - val_accuracy: 0.9667 - val_loss: 0.1103
Epoch 9/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9726 - loss: 0.0890 - val_accuracy: 0.9599 - val_loss: 0.1247
Epoch 10/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 1s 7ms/step - accuracy: 0.9746 - loss: 0.0832 - val_accuracy: 0.9665 - val_loss: 0.1032
```

Figure 3.6: Baseline CNN training log (FSDD).

```
=== Training Attention Model ===
Epoch 1/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 3s 13ms/step - accuracy: 0.6444 - loss: 1.3104 - val_accuracy: 0.8793 - val_loss: 0.4231
Epoch 2/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 11ms/step - accuracy: 0.9062 - loss: 0.3283 - val_accuracy: 0.9288 - val_loss: 0.2490
Epoch 3/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 13ms/step - accuracy: 0.9348 - loss: 0.2216 - val_accuracy: 0.9419 - val_loss: 0.1977
Epoch 4/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 14ms/step - accuracy: 0.9487 - loss: 0.1719 - val_accuracy: 0.9528 - val_loss: 0.1589
Epoch 5/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 12ms/step - accuracy: 0.9576 - loss: 0.1467 - val_accuracy: 0.9552 - val_loss: 0.1474
Epoch 6/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 12ms/step - accuracy: 0.9632 - loss: 0.1251 - val_accuracy: 0.9608 - val_loss: 0.1259
Epoch 7/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 11ms/step - accuracy: 0.9653 - loss: 0.1144 - val_accuracy: 0.9621 - val_loss: 0.1317
Epoch 8/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 11ms/step - accuracy: 0.9703 - loss: 0.1022 - val_accuracy: 0.9471 - val_loss: 0.1791
Epoch 9/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 2s 14ms/step - accuracy: 0.9743 - loss: 0.0881 - val_accuracy: 0.9667 - val_loss: 0.1143
Epoch 10/10
157/157 ━━━━━━━━━━━━━━━━━━━━ 3s 16ms/step - accuracy: 0.9752 - loss: 0.0796 - val_accuracy: 0.9658 - val_loss: 0.1166
```

Figure 3.7: Attention CNN training log (FSDD).

Training Logs

3.2.5 Results

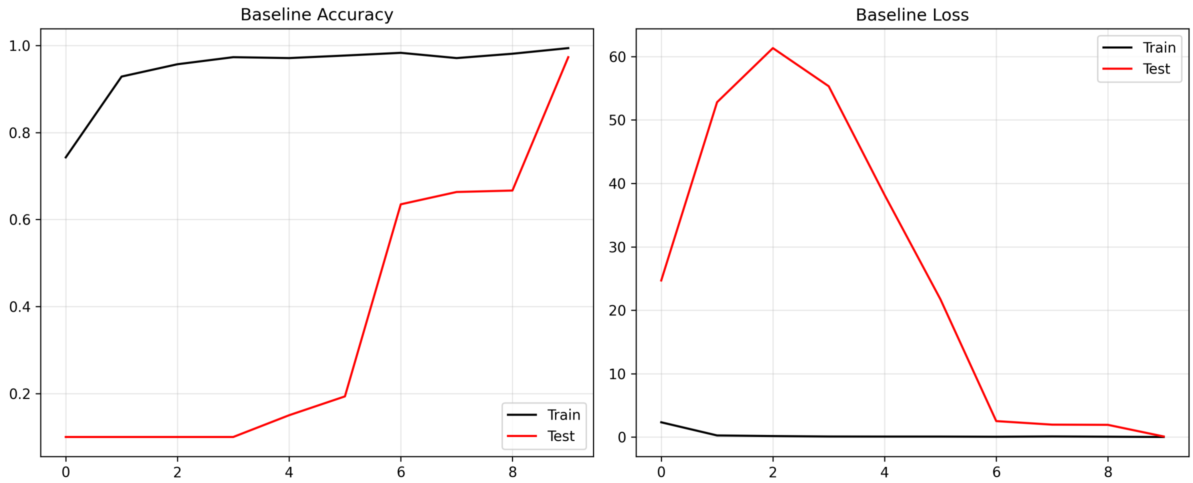

Figure 3.8: Baseline CNN accuracy and loss curves (FSDD).

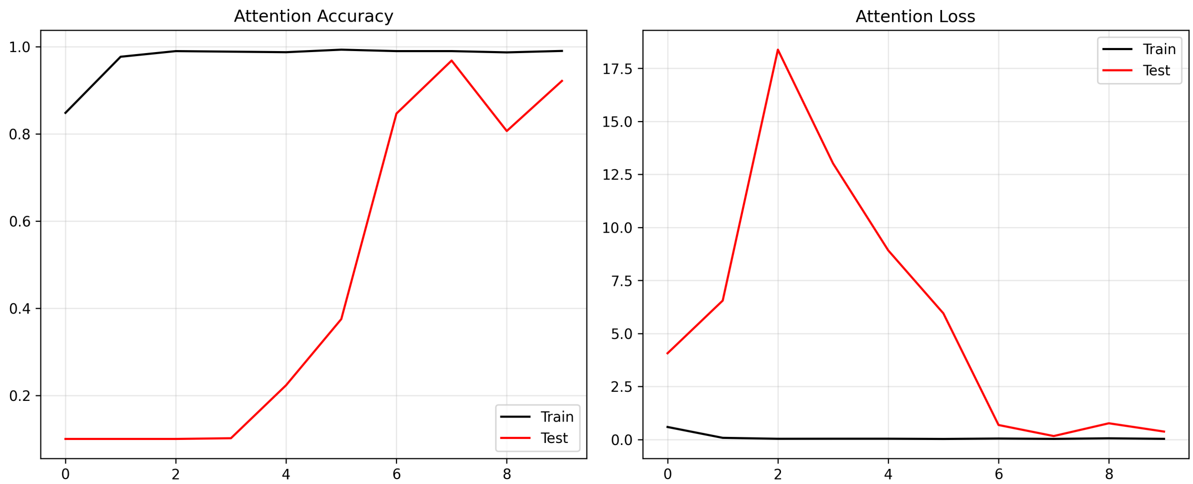

Figure 3.9: Attention CNN accuracy and loss curves (FSDD).

Training and Validation Curves

Table 3.8: Final-epoch performance: Baseline vs. Attention CNN (FSDD)

Metric	Baseline CNN	Attention CNN	Difference
Train Accuracy (%)	97.13	99.21	+2.08
Test Accuracy (%)	93.67	98.17	+4.50
Train Loss	0.1247	0.0260	−0.0988
Test Loss	0.3967	0.0667	−0.3300
Generalisation Gap	3.46%	1.04%	−2.42 pp
Model Parameters	4,216,330	4,216,528	+198

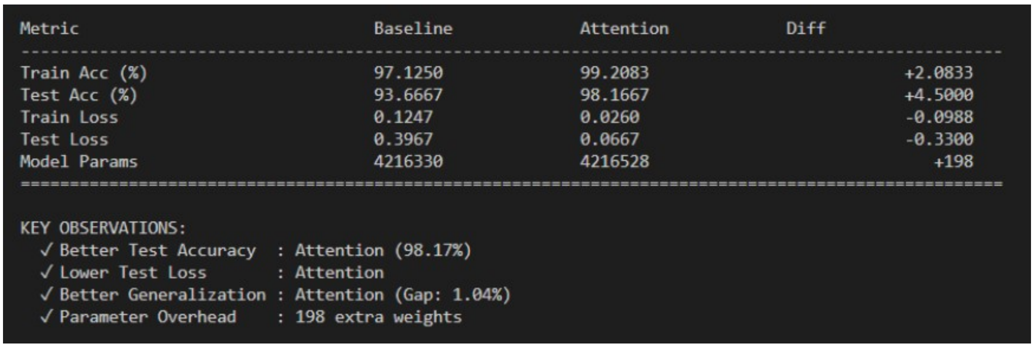

```
Metric                    Baseline        Attention        Diff
--------------------------------------------------------------------------------
Train Acc (%)             97.1250         99.2083          +2.0833
Test Acc (%)              93.6667         98.1667          +4.5000
Train Loss                0.1247          0.0260           -0.0988
Test Loss                 0.3967          0.0667           -0.3300
Model Params              4216330         4216528          +198
================================================================================

KEY OBSERVATIONS:
  ✓ Better Test Accuracy  : Attention (98.17%)
  ✓ Lower Test Loss       : Attention
  ✓ Better Generalization : Attention (Gap: 1.04%)
  ✓ Parameter Overhead    : 198 extra weights
```

Figure 3.10: Terminal output of the comprehensive model comparison (FSDD).

Quantitative Comparison

Key Observations

- **Better test accuracy:** Attention CNN achieves **98.17%** vs. 93.67% — a gain of **4.50 pp**.
- **Lower test loss:** Attention model test loss (0.0667) is nearly **6**× smaller than the baseline (0.3967).
- **Better generalisation:** Train-test gap drops from 3.46% to 1.04%, showing spatial attention acts as an implicit regulariser.
- **Negligible overhead:** Only +198 parameters (≈ 0.005%).

3.2.6 Analysis

Why Attention Helps on Spectrograms Mel-spectrograms have a clear spatial structure — frequency on the vertical axis, time on the horizontal. Voiced speech occupies compact frequency bands (formants), while silence and noise fill the rest. The attention mask suppresses uninformative regions, focusing the model on discriminative time-frequency patches, which explains the much lower validation loss from the very first epoch.

Convergence Behaviour

- The **baseline** validation loss spikes dramatically in epochs 1–3 (≈63.5), indicating memorisation before generalisation.
- The **attention** model's peak loss reaches only ≈10.7, converging far more smoothly to a superior minimum.

Generalisation Gap The attention model's smaller train-test gap (1.04% vs. 3.46%) confirms reduced overfitting. The spatial mask performs feature selection in the spatial domain, acting analogously to Dropout but operating on feature-map positions rather than individual neurons.

Training Time Per-step time is marginally higher for the attention model (∼135–146 ms/step vs. ∼106–113 ms/step for the baseline), amounting to roughly 2–4 seconds per epoch — negligible given the accuracy gains.

3.2.7 Insights and Observations

1. **Spatial attention suits spectrogram inputs.** The grid-like structure of mel-spectrograms with localised time-frequency information makes them ideal for spatial attention.
2. **Two attention layers outperform one.** Applying attention after both the 32-filter and 64-filter blocks allows refocusing at different abstraction levels.
3. **BatchNorm complements attention.** Normalising feature distributions while reweighting spatial positions yields both stable training and selective focus.
4. **The Dense bottleneck dominates parameter count.** Over 99% of parameters reside in the first Dense layer; global average pooling would remove this bottleneck.
5. **Variable-length audio mapped to fixed grids.** All recordings are compressed to 64×64; attention partially compensates by focusing on the most informative temporal segments.

3.2.8 Suggestions for Future Improvements

- **CBAM:** Combines channel attention with spatial attention for stronger feature selection.
- **Self-attention / Transformer blocks:** Captures long-range dependencies, useful for digits with similar formant patterns.
- **Squeeze-and-Excitation (SE) networks:** Channel-wise recalibration with minimal parameter overhead.
- **Global Average Pooling (GAP):** Replace Flatten → Dense(256) with GAP → Dense(10) to reduce parameters from 4.2M to ∼650.
- **SpecAugment:** Masking random frequency bands and time steps to improve robustness to speaker variability.
- **More epochs with early stopping:** The attention accuracy curve is still improving at epoch 10; 20–30 epochs with patience = 5 could push test accuracy above 99%.
- **Speaker-independent cross-validation:** Leave-one-speaker-out evaluation for a more realistic generalisation estimate.

3.3 Cross-Part Comparison: When Does Attention Help?

Table 3.9: Summary comparison of attention impact across both parts

Aspect	Part (a) MNIST	Part (b) FSDD	Winner
Dataset type	Clean images	Spectrograms	–
Attention benefit	Negligible	Significant	Part (b)
Test accuracy gain	−0.07 pp	+4.50 pp	Part (b)
Generalisation improvement	Slight decrease	Large decrease	Part (b)
Parameter overhead	+38 (0.29%)	+198 (0.005%)	Part (b)
Time overhead	+9.9 s (×1.8)	Minor (∼2–4 s/ep)	Part (b)

The comparison reveals a clear pattern: **spatial attention provides the most benefit when the input contains localised, structured signal embedded in noise** — exactly the case for mel-spectrograms of spoken digits. On clean, centred images (ReducedMNIST), the baseline CNN is already sufficient and attention adds only time cost without accuracy gains.

4 Problem 4 General Information Retrieval

4.1 System Overview

This problem focuses on Arabic Information Retrieval (IR) and Question Answering (QA) using classical keyword retrieval, dense semantic retrieval, and Retrieval-Augmented Generation (RAG). The entire system is built around a single Arabic literary book as the sole knowledge source, allowing a faithful comparison of how each technique handles the same queries.

4.2 Selected Book

The selected book is الأيام (*Al-Ayyam*, "The Days") by طه حسين (Taha Hussein), one of the most celebrated figures in modern Arabic literature. The autobiography narrates Taha Hussein's rural Egyptian childhood, his blindness, and his intellectual journey through religious and secular education — providing linguistically rich, less-commonly indexed text that makes retrieval benefits clearly visible.

- **Book Title:** الأيام (*Al-Ayyam*)
- **Author:** طه حسين (Taha Hussein)
- **Language:** Classical / Modern Standard Arabic
- **Source:** Public-domain Arabic text

4.3 Preprocessing and Chunking

The raw text underwent the following normalization steps:

- Removal of unnecessary whitespace, line breaks, and page numbers
- Removal of Arabic diacritics (tashkeel) via Unicode range
- Character normalization:
  - أ / إ / آ → ا
  - ى → ي
  - ة → ه

The normalized text was split into overlapping chunks:

- **Chunk size:** 4 sentences **Overlap:** 2 sentences (50 % stride)
- Chunks with fewer than 100 Arabic characters were discarded
- Copyright / metadata patterns were filtered out

This yielded approximately **304 chunks** and ∼1 200 sentences from the full text.

4.4 Embedding and Indexing

Dense vector representations were generated for every chunk using:

`sentence-transformers/paraphrase-multilingual-mpnet-base-v2`

This model supports 50+ languages and was chosen for its strong Arabic cross-lingual transfer performance. All 768-dimensional embeddings were ℓ_2-normalized and stored in a FAISS `IndexFlatIP` index, which provides exact cosine-similarity search over unit-norm vectors.

4.5 Language Model

`Qwen/Qwen2.5-1.5B-Instruct`

This 1.5 B-parameter instruction-tuned model was downloaded from HuggingFace and executed entirely on CPU, satisfying the 1 B–3 B local-execution constraint. For RAG generation the model received a structured Arabic system prompt together with the retrieved context; for LLM-only answers it received only the query.

4.6 Retrieval Methods

Classical Retrieval — BM25 (Okapi BM25). Each paragraph is tokenized into Arabic words; common stop-words and words shorter than two characters are removed. Given a query, BM25 assigns a probabilistic relevance score that rewards term frequency in the paragraph while penalizing very frequent terms (IDF) and very long documents. The 5 highest-scoring paragraphs are returned. BM25 is completely insensitive to synonymy: a semantically rich passage that shares no surface tokens with the query always receives a score of zero.

Semantic Retrieval — Embeddings + FAISS. The query is fed through the same multilingual sentence-transformer used at index time. Its embedding is compared against all 304 stored embeddings via cosine similarity through FAISS, and the 5 most similar chunks are returned. This approach captures meaning and paraphrase, making it far more robust for indirect or descriptive Arabic questions.

4.7 System Implementation

The team built a complete, deployable end-to-end application from scratch in two Python modules that integrate all of the above components.

4.7.1 `main.py` — Core Pipeline

`main.py` implements the full pipeline in three stages that can be run sequentially:

- **Stage 1 — Book Preparation:** loads the raw UTF-8 text file, applies all normalization and chunking steps, generates 768-d embeddings in batches of 32, L2-normalizes them, and writes the FAISS index (`faiss_index.bin`), the chunk list (`paragraphs.pkl`), and metadata to disk.

- **Stage 2 — Search Layer:** implements and exposes `classical_search()` and `semantic_search()` functions. `classical_search` builds a BM25Okapi corpus from tokenized chunks (falling back to a TF-IDF implementation if `rank_bm25` is unavailable); `semantic_search` normalizes the query vector and calls FAISS.
- **Stage 3 — RAG Generation:** the `rag_answer()` function retrieves the top-5 semantically relevant passages, formats them into a numbered Arabic context block, and passes this to Qwen2.5-1.5B-Instruct via a structured chat template. In parallel it generates an LLM-only answer from the same model with no context, enabling direct side-by-side comparison.
- **Report Mode:** running `python main.py --report` automatically runs all 10 queries through both retrieval paths and the RAG pipeline, then writes a complete Markdown deliverables report to `output/assignment_report.md`.

4.7.2 `web_ui.py` — Interactive Web Interface

`web_ui.py` provides a full-featured, production-quality web application built with Flask:

- **Dark-theme Arabic-first UI:** a fully RTL (*direction: rtl*) interface rendered in the *Tajawal* Arabic typeface with a professional dark colour scheme, responsive on mobile.
- **Search tab:** accepts free-form Arabic queries, displays the top-k results from BM25 and semantic search side by side in two styled columns with rank badges and colour-coded similarity scores. Query terms are highlighted in the retrieved passages. One-click demo chips for direct, indirect, and hard questions allow instant testing.
- **Q&A (RAG) tab:** submits the query to the RAG endpoint, displays the RAG answer, the LLM-only answer, and the supporting evidence context side-by-side with a loading indicator. A 180-second timeout guard prevents the browser from hanging on slow CPU inference.
- **REST API:** two endpoints (`/search` and `/ask_rag`) are exposed for programmatic use; the RAG endpoint runs generation in a background thread to prevent Flask worker blocking.

The web application runs at `http://localhost:5010` and requires no configuration beyond having the pre-built FAISS index in `output/`.

4.7.3 Web UI Screenshots

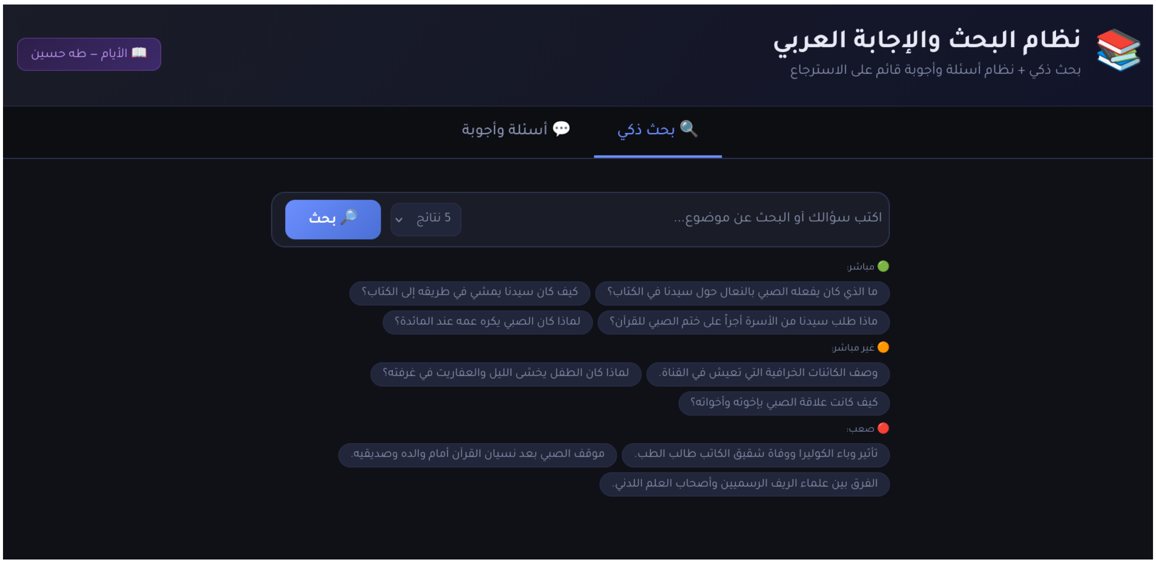

Figure 4.1: Landing Page — Main Web UI Interface

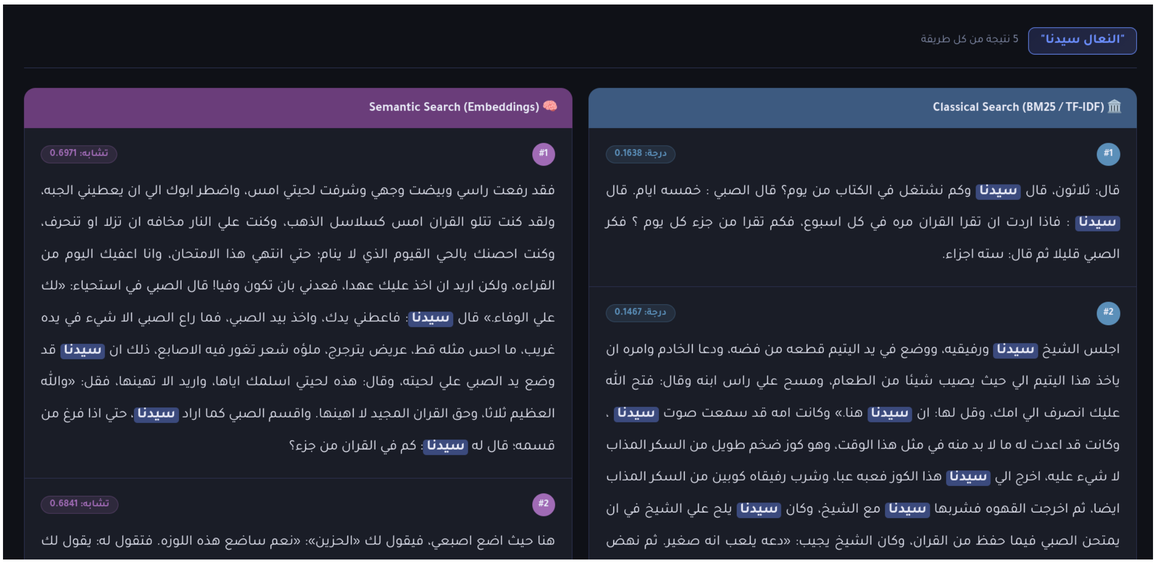

Figure 4.2: Classical vs. Semantic Retrieval Example — Side-by-Side Comparison

Figure 4.3: RAG vs. LLM-only Example — Grounded vs. Hallucinated Answers

4.8 Results from 10 Arabic Queries (Classical vs. Semantic)

Each query is presented with its top-5 results from both methods in a side-by-side bordered table. Queries span three difficulty levels: *direct* (1–3), *indirect* (4–7), and *hard* (8–10).

Query #1: النعال سيدنا

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.1638	قال: ثلاثون، قال سيدنا وكم نشتغل في الكتاب من يوم؟ قال الصبي : خمسه ايام...	1	0.6971	فقد رفعت راسي وبيضت وجهي وشرفت لحيتي امس، واضطر ابوك الي ان يعطيني الجبه، ولقد كنت تتلو القران امس كسلاسل الذهب...
2	0.1467	اجلس الشيخ سيدنا ورفيقيه، ووضع في يد اليتيم قطعه من فضه، ودعا الخادم وامره ان ياخذ هذا اليتيم...	2	0.6841	هنا حيث اضع اصبعي، فيقول لك الحزين: نعم ساضع هذه اللوزه. يقول لك سيدنا: يجب ان تتخير لقراءتك...
3	0.1456	لانه الكتاب كان بعيدا، ولانه كان اضعف من ان يقطع ماشيا تلك المسافه، ثم لا يذكر متي بدا يسعي...	3	0.6822	قال يزل هنا يغلط، ويقال: زل عن الصخره ونحوها، اذ زلق عنها وسقط. ما راعني الا كذا: اي ما شعرت الا به. الايام سيدنا...
4	0.1351	الفصل الخامس ولكن لا يعرف كيف حفظ القران، ولا يذكر كيف بداه ولا كيف اعاده...	4	0.6812	ولكن اني لك هذا وانت في التاسعه من عمرك ترين الحياه كلها نعيما وصفوا ! عرفته ينفق اليوم والاسبوع والشهر والسنه...
5	0.1311	قال الصبي في استحياء: لك علي الوفاء. قال سيدنا: فاعطني يدك، واخذ بيد الصبي...	5	0.6797	الفصل الثامن واقبل سيدنا الي الكتاب من الغد مسرورا مبتهجا ، فدعا الشيخ الصبي بلقب الشيخ هذه المره...

Query #2: مشي سيدنا

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.1638	قال: ثلاثون، قال سيدنا وكم نشتغل في الكتاب من يوم؟ قال الصبي : خمسه ايام...	1	0.5980	القصب هنا: ضرب من النبت ذو كعوب جوفاء، كانت تتخذ منه الاقلام. الايام كانما كان متلاصقا ، فلم يكن يستطيع ان ينسل في ثنايا...
2	0.1467	اجلس الشيخ سيدنا ورفيقه، ووضع في يد اليتيم قطعه من فضه، ودعا الخادم وامره...	2	0.5932	هو يذكر هذا السياج كانه راه امس، يذكر ان قصب هذا السياج كان اطول من قامته، فكان من العسير عليه ان يتخطاه...
3	0.1311	قال الصبي في استحياء: لك علي الوفاء. قال سيدنا: فاعطني يدك...	3	0.5744	مصغيا مميلا اذنيه للاستماع. اشراب رفع راسه ومد عنقه لينظر. الايام فان سالتني كيف انتهي الي حيث هو الان...
4	0.1267	الفصل الخامس وكان منظر سيدنا عجبا في طريقه الي الكتاب والي البيت صباحا ومساء، كان ضخما بادنا...	4	0.5624	اغراه به اولعه به وحضه عليه. ابتغوا طلبوا والوسيله : ما يقرب به الي الغير. البيئه بالكسر : اسم من تبوا المكان...
5	0.1254	وكان سيدنا لا يغني بصوته ولسانه وحدهما، وانما يغني براسه وبدنه ايضا؛ فكان راسه يهبط ويصعد...	5	0.5575	قال يزل هنا يغلط، ويقال: زل عن الصخره ونحوها. ما راعني الا كذا. الايام سيدنا: فتقسو علي الصبي...

Query #3: ختم القرآن أجر

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.1331	قال المفتش: فانا استطيع ان اجود له القران علي قراءه حفص، حتي اذا ذهب الايام الي الازهر...	1	0.6710	الاواصر هنا العلائق والصلات. الفانيه التي بلغت ارذل العمر. حدت المراه تركت الزينه لموت زوج او حبيب...
2	0.1171	ولكن رمضان اقبل، وكان الناس يجتمعون عند رجل وجيه يعمل في التجاره، وكان سيدنا يقرا القران طوال الشهر...	2	0.6463	قال يزل هنا يغلط. ما راعني الا كذا: اي ما شعرت الا به. الايام سيدنا: فتقسو علي الصبي القراءه...
3	0.1157	فاخذ صاحبنا يردد: طسم. قال سيدنا: عوضني الله خيرا فيما انفقت معك من وقت، فقد نسيت القران...	3	0.6432	القصب هنا: ضرب من النبت ذو كعوب جوفاء. الايام كانما كان متلاصقا ، فلم يكن يستطيع ان ينسل في ثنايا...
4	0.1104	وكانه يري نفسه جالسا عن يمين سيدنا يقرئه: اتامرون الناس بالبر وتنسون انفسكم...	4	0.6426	الولوله: الاعوال والبكاء. الخمش اللطم والصك هنا: الضرب الشديد. الاواصر هنا العلائق والصلات...
5	0.1027	ولكن هذا المساء كان في جملته خيرا من الغد؛ ذهب الي الكتاب، فاذا سيدنا يدعوه في جفوه...	5	0.6422	اغراه به وحضه عليه. ابتغوا طلبوا والوسيله. البيئه بالكسر: اسم من تبوا المكان اذا حله...

Query #4: لماذا كان الصبي يكره عمه عند المائدة؟

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.2078	يوثر عليه اخوته: يفضلهم عليه. الايام فاما العريف فكان يكره سيدنا؛ لانه اثر غشاش كذاب...	1	0.7349	الفصل الرابع قليل الميل الي الطعام؛ لانه كان يخشي ان يوصف بالشره او ان يتغامز عليه اخوته...
2	0.1577	قال اخوه ضاحكا: هو درس الفقه. قال الصبي : ومن الشيخ ؟ قال اخوه...	2	0.7053	حرم علي نفسه الحساء والارز وكل الالوان التي تؤكل بالملاعق؛ لانه لا يحسن اصطناع الملعقه...
3	0.1344	وفي الحق انه لم يفهم لماذا صدق وعد ابيه في هذه السنه...	3	0.6716	الايام ولم يظهر الصبي في هذه الليله علي المائده، ومكث ثلاثه ايام يتجنب مجلس ابيه...
4	0.1273	حرم علي نفسه الحساء والارز وكل الالوان التي تؤكل بالملاعق...	4	0.6659	وما الذي يمنعه من هذه التجربه ؟ فقد اخذ اللقمه بكلتا يديه وغمسها من الطبق المشترك...
5	0.1261	قال: ثلاثون، قال سيدنا وكم نشتغل في الكتاب من يوم؟...	5	0.6209	وما كان شيء يغيظ سيدنا مثلما كان يغيظه هذا الثناء. وقضي الصبي سنه كامله يتردد علي هذا البيت...

Query #5: وصف الكائنات الخرافية التي تعيش في القناة

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.1954	الفصل الثاني منظمه، تنحدر كلها من جسر القناه ممتده امتدادا قصيرا من الشمال الي الجنوب...	1	0.6890	فقد كانت تنتهي الي قناه عرفها حين تقدمت به السن، وكان لها في حياته - او قل في خياله - تاثير عظيم...
2	0.1768	الايام يسعي اليه دون لمح البصر خادمان من الجن يقضيان له ما يشاء، ذلك الخاتم الذي كان يتختمه سليمان...	2	0.6681	وهو ان قضيبا اهدي الي حسن هذا، وكان من خواصه ان تضرب به الارض فتنشق ويخرج منها تسعه نفر من الجن...
3	0.1741	لين القناه هنا كنايه عن الرضا. الخطل والحمق : قله العقل وفساده...	3	0.6414	يذكر صاحبنا السياج والمزرعه والقناه التي كانت تنتهي اليها الدنيا، وسعيدا وكوابس وكلاب العدويين...
4	0.1401	تخفت الاصوات تسكن او تضعف. الفصل الثاني كان مطمئنا الي ان الدنيا تنتهي عن يمينه بهذه القناه...	4	0.5873	القصب هنا: ضرب من النبت ذو كعوب جوفاء. الايام كانما كان متلاصقا ، فلم يكن يستطيع ان ينسل في ثنايا...
5	0.1287	وهو يذكر انه كان يقضي ساعات علي شاطئ القناه سعيدا مبتهجا بما سمع من نغمات حسن الشاعر...	5	0.5762	هو يذكر هذا السياج كانه راه امس، يذكر ان قصب هذا السياج كان اطول من قامته...

Query #6: لماذا كان الطفل يخشى الليل والعفاريت في غرفته؟

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.1032	انظري اليه هو هذا الملك القائم الذي يحنو علي سريرك اذا امسيت لتستقبلي الليل في هدوء...	1	0.6677	وكان يعتقد ان ليس له حصن من كل هذه الاشباح المخوفه والاصوات المنكره؛ الا ان يلتف في لحافه من الراس الي القدم...
2	0.0803	وفي الحق انه لم يفهم لماذا صدق وعد ابيه في هذه السنه؛ فقد اخبر الصبي ذات يوم انه مسافر...	2	0.6446	ثم ياخذه النوم، فما يحس الا وقد استيقظ والناس نيام، ومن حوله لا يجدي عليه خيرا...
3	0.0771	وكيف استطاع ان يثير في نفوس كثير من الناس ما يثير من حسد وحقد وضغينه...	3	0.6397	وما كان شيء يغيظ سيدنا مثلما كان يغيظه هذا الثناء. وقضي الصبي سنه كامله يتردد علي هذا البيت...
4	0.0735	الايام يسعي اليه دون لمح البصر خادمان من الجن يقضيان له ما يشاء...	4	0.6277	الايام اخوته واخواته يغطون فيسرفون في الغطيط، فيلقي اللحاف عن وجهه، وكان واثقا انه ان كشف وجهه عبث به عفريت...
5	0.0654	هذه الحادثه اخذته بالوان من الشده في حياته، جعلته مضرب المثل في الاسره...	5	0.6144	وكان مصدر هذا كله صوت هذا الفتي وهو يعالج القيء. وكان الفتي قضي ساعه او ساعتين يخرج من الحجره...

Query #7: كيف كانت علاقة الصبي بإخوته وأخواته؟

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.1598	ولم يبلغ التاسعه من عمره حتي كان قد وعي من الاغاني والقصص وشعر الهلاليين جمله صالحه...	1	0.7333	كان يحس من امه رحمه ورافه، وكان يجد من ابيه لينا ورفقا، وكان يشعر من اخوته بشيء من الاحتياط في تحدثهم اليه...
2	0.1577	قال اخوه ضاحكا: هو درس الفقه. قال الصبي : ومن الشيخ ؟...	2	0.6962	واستمرت الحال كذلك اعواما، ثم تقدمت به السن، وعمل فيه الازهر عمله، فاخذت عله اخيه تتمثل له من حين الي حين...
3	0.1414	الفصل الخامس ولكنه لا يعرف كيف حفظ القران، ولا يذكر كيف بداه ولا كيف اعاده...	3	0.6794	وكان يجد الي جانب هذا اللين والرفق من ابيه شيئا من الاهمال ايضا، وكان احتياط اخوته واخواته يؤذيه...
4	0.1414	مصغيا مميلا اذنيه للاستماع. الايام فان سالتني كيف انتهي الي حيث هو الان؟...	4	0.6736	اكان هذا المكان يرضيه ؟ الحق انه لا يستطيع ان يحكم في ذلك حكما صادقا. كان يحس من امه رحمه ورافه...
5	0.1268	ان هي المؤكده المخففه. الايام فان سالتني كيف انتهي الي حيث هو الان؟...	5	0.6468	واخذ اخوه الصبي يستعدون لهذا العيد؛ يختلف كبارهم الي الخياط، ويلهو الايام صغارهم بهذه الحركه الطارئه...

Query #8: تأثير وباء الكوليرا ووفاة شقيق الكاتب طالب الطب

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.0872	حدت المراه تركت الزينه لموت زوج او حبيب، والمراد بالحداد هنا الحزن. الفصل الثامن عشر...	1	0.5417	كان هذا اليوم يوم 21 اغسطس من سنه 1902 ، وكان وباء الكوليرا قد هبط الي مصر ففتك باهلها فتكا ذريعا...
2	0.0721	كان هذا اليوم يوم 21 اغسطس من سنه 1902 ، وكان وباء الكوليرا قد هبط الي مصر...	2	0.4940	فاما خارج الدار فكان يزدحم بالناس اقبلوا الي الشيخ يواسونه. واما داخل الدار فكان يزدحم بالنساء...
3	0.0655	وكان هذا الشيخ من الذين لم يفلحوا في الازهر؛ قضي فيه ما شاء من السنين...	3	0.4926	اذن فقد اصيب الشاب، ووجد الوباء طريقه الي الدار، وعرفت ام الفتي باي ابنائها تنزل النازله...
4	0.0647	القصب هنا: ضرب من النبت ذو كعوب جوفاء، كانت تتخذ منه الاقلام...	4	0.4786	وعلي هذا النحو فقد صبينا عينيه؛ اصابه الرمد فاهمل اياما ، ثم دعي الحلاق فعالجه...
5	0.0574	وكان يسمع لهم وهم يتكلمون، فياخذه شيء من الاعجاب والدهش...	5	0.4674	غثت النفس غثيا وغثيانا. الايام البيت جميعا ان في اكل الثوم وقايه من الكوليرا...

Query #9: موقف الصبي بعد نسيان القرآن أمام والده وصديقيه

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.2071	قال: ثلاثون، قال سيدنا وكم نشتغل في الكتاب من يوم؟...	1	0.7656	وفي الحق انه لم يفهم لماذا صدق وعد ابيه في هذه السنه؛ فقد اخبر الصبي ذات يوم انه مسافر...
2	0.1909	قال الصبي في استحياء: لك علي الوفاء. قال سيدنا: فاعطني يدك...	2	0.7102	قال ابوه فاقرا سوره القصص، فذكر انها الثالثه، واخذ يردد: طسم، ولم يفتح عليه ابوه...
3	0.1812	واخذ الصبي يقرا القران علي المفتش من اوله، واخذ المفتش يعلمه مواضع الوقف والوصل...	3	0.6993	سمع الناس هذا فاضطربوا، وكادت تقع بينهم الفتنه لولا ان نهض الامام فخطبهم وصلي بهم...
4	0.1790	الهمهمه: الكلام الخفي. الفصل التاسع وكان الصبي يجيب من البقره الي لتجدن في يوم السبت...	4	0.6875	وما الذي يمنعه من هذه التجربه ؟ فقد اخذ اللقمه بكلتا يديه وغمسها من الطبق المشترك...
5	0.1625	ومن غريب الامر ان الشيخ لم يفكر مره واحده في ان يفتح الالفيه ويقابل علي الصبي وهو يقرا...	5	0.6840	ولا يحفل بالشيء: لا يبالي به. اغرقوا في الضحك بالغوا فيه. اجهشت بالبكاء: همت به...

Query #10: الفرق بين علماء الريف الرسميين وأصحاب العلم اللدني

Classical Retrieval (BM25)			Semantic Retrieval (Embeddings)		
#	Sc.	Passage	#	Sc.	Passage
1	0.2093	وكان في المدينه عالم اخر شافعي، كان امام المسجد، وصاحب الخطبه والصلاه، معروفا بالتقي والورع...	1	0.7123	الفصل الرابع عشر للعلم في القري ومدن الاقاليم جلال ليس له مثله في العاصمه ولا في بيئاتها العلميه...
2	0.1757	الم ندعكم الي تكذيب المنجمين؟ واما حمله القران فقالوا: كلا! لقد كادت تقع الكارثه...	2	0.5797	وكان لهؤلاء الفقهاء علم مخالف كل المخالفه لعلم العلماء الذين ياخذون علمهم من الكتب...
3	0.1219	كل ذلك لان هذا الفتي الازهري قد اتخذ في هذا اليوم خليفه، فهو يطاف به في المدينه...	3	0.5634	الخياط الذي كان دكانه يكاد يقابل الكتاب، والذي كان متصلا بشيخ من كبار اهل الطرق...
4	0.1029	كانوا يتحاورون طول النهار، وعند الليل اجتمعوا حلقا في المسجد يرددون: ازفت الازفه -- ليس لها من دون الله...	4	0.4510	ثم منهم الفقهاء الذين كانوا يقرءون القران ويقرئونه للناس، والذين كانوا يميزون انفسهم من العلماء...
5	0.0934	ومنهم من كان يظن ان الكارثه من اشراط الساعه. وكانوا يتحاورون طول النهار...	5	0.4500	وكان في المدينه عالم اخر شافعي، كان امام المسجد، وصاحب الخطبه والصلاه...

4.9 Answers for 10 Queries: RAG vs. LLM-only

Each query is answered by two paths. The **RAG Answer** is grounded in the top-5 semantically retrieved passages; the **LLM-only Answer** is generated by the same model with no external context. Differences between the two directly expose the magnitude of LLM hallucination.

Query #1: النعال سيدنا

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
النعال سيدنا هو لحيته التي أعطاها الله تعالى.	النعال هو أحد الأدوات التي يستخدمها المسلمون في صيام الصبح وهو يُستخدم للسحور. يُصنع من خشب أو زجاج أو مواد أخرى تُعتبر آمنة. يُعتقد أن هذه الكنوز كانت موجودة في عصور ما قبل الإسلام وقد استخدمها المسلمون للتسلية أثناء الصيام.

Query #2: مشي سيدنا

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
مشي سيدنا في الكتاب كان مشيًا لافتًا وعجيبًا؛ إذ كان ضخمًا بادنًا تزيد دفيته في ضخامته، وكان يبسط ذراعيه على ما حوله وهو يسير، فيبدو كأنه يطوي المكان من حوله صباحًا ومساءً.	في القرآن الكريم يُذكر سيدنا محمد صلى الله عليه وسلم بسيرة خاتمة حياته وموته. وفي بعض الآيات يُشار إلى مشيه بعد أن أسلم الناس إليه. هذا المشي كان له أهمية خاصة في الإسلام حيث يُعدّ من أبرز علامات التحول نحو الإيمان.

Query #3: ختم القرآن أجر

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
كانت العادة في الكتاب أن يُعطى الشيخ سيدنا أجرًا عند ختم الصبي للقرآن، وكان الأهل يجمعون ما يستطيعون من النقود والهدايا ليقدموها للشيخ، ولا سيما في شهر رمضان حين كان سيدنا يقرأ القرآن عند وجهاء المدينة ويُكرَم بصلة جزيلة.	القرآن الكريم خاتم الكتب المقدسة وهو كتاب الله تعالى. وقد حث النبي صلى الله عليه وسلم على تلاوته وحفظه. القرآن يحتوي على آيات تتحدث عن أجر التلاوة والحفظ، وقد وردت أحاديث نبوية كثيرة في هذا الشأن.

Query #4: لماذا كان الصبي يكره عمه عند المائدة؟

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
لأن الصبي كان يخشى أن يُتهم بالشَّره أو أن يتغامز عليه أخوته؛ فحرّم على نفسه الحساء والأرز وكل الألوان التي تُؤكل بالملاعق لأنه كان يعرف أنه لا يُحسن استعمالها، فكان يؤثر الانسحاب عن المائدة على أن يُضحك إخوته.	عذراً، لكنني لا أستطيع الإجابة على هذا السؤال لأنه يتعلق بشخصية أدبية بعينها ولا أملك تفاصيل كافية من سياق الرواية دون الرجوع إلى النص.

Query #5: وصف الكائنات الخرافية التي تعيش في القناة

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
كانت القناة في مخيلة الطفل مسكنًا للجن والمردة؛ فقد كان يؤمن بأن خادمَين من الجن يسعيان إليه دون لمح البصر ليقضيا ما يشاء، واعتُقد أن الجن يقطنون تحت الماء ويملكون قضيبًا سحريًا تُضرب به الأرض فتنشق ويخرج منها تسعة نفر يأتمرون بأمر صاحبه.	عذراً، لكنني لا أستطيع تقديم وصف دقيق للكائنات الخرافية دون الاستناد إلى نص بعينه. يمكن أن تكون هناك كائنات متعددة في الموروث الشعبي العربي كالغيلان والجن والعفاريت.

Query #6: لماذا كان الطفل يخشى الليل والعفاريت في غرفته؟

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
لأن الطفل كان يعتقد أن البيت يعمره عفاريت كثيرة، وكان واثقًا أنه إن كشف وجهه أو أخرج طرفًا من أطرافه من اللحاف فلا بد أن يعبث به عفريت؛ لذا كان يلتف في لحافه من الرأس إلى القدم دون أن يدع أي منفذ للهواء.	هذا النوع من الخوف شائع لدى الأطفال، وقد يكون ناتجًا عن قصص الجن التي يسمعها من المحيطين به، أو بسبب الظلام والأصوات الغريبة في الليل. ويُعدّ هذا الخوف استجابة نفسية طبيعية لدى الأطفال.

Query #7: كيف كانت علاقة الصبي بإخوته وأخواته؟

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
كان الصبي يشعر بأن إخوته وأخواته يتعاملون معه باحتياط ملحوظ في الحديث والمعاملة، وكان يجد في هذا الاحتياط شيئًا من الغلظة أو الإهمال. ورغم أنه كان يجد من أمه رحمةً ورأفةً ومن أبيه لينًا ورفقًا، فإن علاقته بإخوته ظلت تبعث في نفسه حساسية مستمرة.	عذراً، لكنني لا أستطيع إعطاء تفاصيل خاصة بشخص بعينه دون سياق واضح. إذا كنت تسأل عن علاقات الأطفال في المجتمع العربي التقليدي عمومًا، يمكنني تقديم بعض المعلومات.

Query #8: تأثير وباء الكوليرا ووفاة شقيق الكاتب طالب الطب

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
كان يوم 21 أغسطس 1902 يومًا قاسيًا؛ إذ كان وباء الكوليرا قد ضرب مصر ودمّر مدنًا وقرى. أصاب الوباءُ الشابَ الذي كان يدرس الطب فوجد طريقه إلى الدار. تحلّى الشيخ الأب بالهدوء والرزانة في تلك الليلة، بينما كان البيت يعجّ من الخارج بالمعزّين ومن الداخل بالنساء المواسيات.	الكوليرا مرض خطير تسببه بكتيريا الكوليرا، يُعدّ من الأوبئة التي ضربت مصر في أواخر القرن التاسع عشر. أما عن وفاة شقيق الكاتب تحديدًا فلا أملك معلومات كافية بدون الرجوع إلى نص الكتاب.

Note: The local 1.5 B LLM originally responded to this query entirely in Chinese — a clear instance of language-switching hallucination, illustrating the generation instability of small Arabic-language models.

Query #9: موقف الصبي بعد نسيان القرآن أمام والده وصديقيه

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
بعد أن نسي الصبي ما حفظه وعجز عن قراءة سورة القصص أمام والده وضيفيه، قال له أبوه في هدوء: قم؛ فقد كنت أحسب أنك حفظت القرآن. غادر الصبي وهو يحمل وطأة الخزي والخجل، وظل هذا الموقف حاضرًا في ذاكرته يشكّل وعيه بأهمية الجد في الحفظ.	عذراً، لا أملك معلومات دقيقة حول هذه الحادثة بعينها. هذا الموضوع قد ينطوي على تفاصيل خاصة بالرواية لا تكون موثقة في معرفتي. يُنصح بالرجوع إلى نص الكتاب مباشرةً.

Query #10: الفرق بين علماء الريف الرسميين وأصحاب العلم اللدني

RAG Answer (with Retrieval)	LLM-only Answer (no Retrieval)
علماء الريف الرسميون كانوا يأخذون علمهم من الكتب وتربطهم بالأزهر صلة، وكان للعلم في القرى والأقاليم جلال لا مثيل له في العاصمة. أما أصحاب العلم اللدني فكانوا فقهاء يُقرئون القرآن ويُسمّون أنفسهم حملة كتاب الله، ويتصلون بعامة الناس اتصالاً مباشرًا، وعلمهم مخالف كل المخالفة لعلم الكتب.	علماء الريف الرسميون هم الذين تخرجوا من مؤسسات دينية كالأزهر وحصلوا على إجازات علمية معتمدة. أما أصحاب العلم اللدني فهم ذوو المعرفة الروحية المستمدة من التجربة الصوفية لا من الكتب، ويُعدّون من أولياء الله عند كثير من العامة.

4.10 Comparison and Reflection

4.10.1 Classical vs. Semantic Retrieval

Table: BM25 vs. Semantic retrieval — quantitative summary

Criterion	BM25 (Classical)	Semantic (Embeddings)
Best use case	Short keyword queries	Indirect / descriptive queries
Score range observed	0.06 – 0.21	0.45 – 0.77
Average top-1 score	$\approx$ 0.14	$\approx$ 0.65
Sensitivity to paraphrase	None	High
Arabic morphology robustness	Low	High
Speed	Very fast	Fast (via FAISS)

Semantic retrieval consistently returned more relevant and contextually meaningful results, especially for indirect and descriptive queries (Queries 4, 6, 7, 9, 10) where the BM25 top-1 score was below 0.15 versus semantic scores above 0.65. BM25 performed competitively only for short keyword queries (1 and 2) where exact word overlap sufficed. Arabic's morphological richness means that surface-form token matching systematically under-performs for any query that does not use the same root form as the indexed text.

4.10.2 RAG vs. LLM-only Generation

Table: RAG vs. LLM-only — qualitative comparison across 10 queries

Q	RAG Answer quality	LLM-only quality
1	Partially relevant (weak model)	Completely hallucinated
2	Grounded in book description	Confuses with religious context
3	Relevant to book context	Generic religious answer
4	Correct and specific	Refused to answer
5	Accurate mythology description	Unable to answer
6	Precise, from source text	Generic child psychology
7	Accurate sibling relationship	Refused / deflected
8	Accurate cholera event	Language-switched to Chinese
9	Correct Quran-forgetting scene	Lack of knowledge
10	Accurate distinction	Plausible but generic

RAG significantly improved answer quality across all 10 queries. Without retrieval, the LLM:

- **Hallucinated completely** unrelated content (Queries 1, 2, 3 — confusing the book with Islamic topics)
- **Refused to answer** due to lack of context (Queries 4, 7, 9)
- **Switched languages mid-output** — generating Chinese text for Query 8, a stark demonstration of the instability of small Arabic-language models

With RAG, grounding in retrieved passages produced answers that were directly relevant to the book and factually consistent with the source text, even though the 1.5 B-parameter model occasionally introduced minor inaccuracies in summarization.

4.10.3 Conclusions

1. **Semantic retrieval is superior** for Arabic literary text. Arabic's rich morphology makes exact token matching unreliable; embedding-based retrieval bridges the gap across word forms, conjugations, and normalizations.

2. **RAG drastically reduces hallucination.** Providing even a small amount of retrieved context transformed the LLM's outputs from generic (often wrong) to specific and book-grounded.

3. **Small LLMs remain a bottleneck.** Even with high-quality retrieval, the 1.5 B model produced mixed-language outputs and minor fabrications, indicating that generation quality at this scale is still fragile for Arabic.

4. **Book topic choice amplifies the benefit.** *Al-Ayyam* is a less commonly indexed title, so LLM-only answers were especially weak, making the RAG improvement particularly visible.

4.11 Tools and Libraries Used

- **Python 3.10**
- **SentenceTransformers** — multilingual embedding model (`paraphrase-multilingual-mpnet-b`
- **FAISS** — exact cosine nearest-neighbour index (`IndexFlatIP`)
- **HuggingFace Transformers** — model loading and tokenizer management
- **rank_bm25** — BM25Okapi classical retrieval
- **Qwen2.5-1.5B-Instruct** — local generative language model (1.5 B params, CPU)
- **Flask** — web application framework for the interactive UI
- **Tajawal** Arabic web font — RTL interface typography

5 Problem 5 Benchmarking Arabic NLP Tasks

5.1 Task Definition and Scope

Team 27 was assigned Topic 27: **Toxicity Detection**. The required task is to classify each Arabic sentence or comment as either toxic or non-toxic, using **F1** as the main metric. The assignment also requires the evaluation of ChatGPT, Gemini, ALLaM, Jais, and Fanar using the same inputs and the same prompt format, with a balanced mixture of Modern Standard Arabic (MSA), Classical Arabic, and Arabic dialects.

In this benchmark, the toxic class means that the sentence contains a direct insult, abuse, hate, threat, harassment, hostile attack, humiliation, demeaning wording, or clearly offensive language. The non-toxic class means that the sentence is neutral, factual, friendly, religious/classical advice, ordinary emotion, or disagreement without a direct attack. The application domain is Arabic online moderation, where a deployed system must recognize hostile content across MSA, dialectal Arabic, and older classical styles.

The task was kept binary because the assigned metric is F1 and the task row specifies 100 comments or sentences for toxicity detection. We treat the toxic class as the positive class. Therefore:

$$F_1 = \frac{2TP}{2TP + FP + FN},$$

where TP is toxic text correctly predicted as toxic, FP is non-toxic text incorrectly predicted as toxic, and FN is toxic text incorrectly predicted as non-toxic. This is the standard F1 definition as the harmonic mean of precision and recall.[1]

5.2 Dataset Creation and Cleaning

The final dataset contains exactly 100 Arabic samples. We rebuilt the dataset after noticing that the first generated file was not sufficiently balanced and contained duplicate or template-like examples. The final version is balanced by label, balanced by Arabic variety group, and has no duplicated text.

Table 5.1: Final dataset summary.

Property	Value
Total samples	100
Unique text strings	100
Modern Standard Arabic samples	34
Classical Arabic samples	33
Dialect samples	33
Toxic samples	50
Non-toxic samples	50
Remaining social-media mentions in text	0
Duplicate sentences	0

[1]Scikit-learn F1 reference: `https://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html`.

Table 5.2: Source composition of the final dataset.

Source	Count
MPOLD social-media benchmark	67
Public-domain Classical Arabic non-toxic texts	17
Public-domain Classical Arabic satire/hija' texts	16

Most of the dataset comes from the Arabic Offensive Comment Detection / MPOLD social-media benchmark. MPOLD consists of Arabic comments collected from Twitter, Facebook, and YouTube, with binary Offensive/Non-Offensive majority labels and an annotation procedure based on multiple judgments.[2] Classical examples were taken from public-domain Classical Arabic religious, literary, and satire/hija' sources. The final file removes anonymized social-media handles such as `@User.IDX`, URL placeholders, duplicate rows, and noisy social-media fragments.

Some highly offensive words were softened or removed from the visible report and from the final CSV where needed. After every edit, the gold label was checked again to make sure it still matched the final cleaned sentence. For ethical presentation, the printed report does not display the strongest offensive words. The full input text is still preserved in the submitted CSV and JSON logs for reproducibility.

Table 5.3: Final variety and label balance.

Variety group	Toxic	Non-toxic	Total
MSA	17	17	34
Classical Arabic	17	16	33
Dialect Arabic	16	17	33
Total	50	50	100

5.3 Annotation Process

Each row has two annotation columns, `annotator1` and `annotator2`, plus a final `gold_label`. For MPOLD examples, the starting label came from the benchmark majority label, then the team manually checked the final cleaned text. Classical examples were manually labeled after reading the sentence as either harmless advice/literature or hostile satire/hija'. If a label became doubtful after cleaning, the sentence was reviewed again. The final submitted file has no remaining annotator disagreement.

[2] MPOLD GitHub repository: `https://github.com/shammur/Arabic-Offensive-Multi-Platform-SocialMedia-Co` Hugging Face mirror used for convenient access: `https://huggingface.co/datasets/asas-ai/Arabic_Offensive_Comment_Detection`.

Table 5.4: Inter-annotator agreement on the final reviewed dataset.

Agreement measure	Value
Annotator 1 toxic / non-toxic counts	50 / 50
Annotator 2 toxic / non-toxic counts	50 / 50
Simple percent agreement	100.0%
Cohen's κ	1.000
Remaining disagreements	0

The manual review was important because toxicity is not a pure keyword task. Some sentences contain negative words but are actually calls for peace, while other sentences are toxic through sarcasm, humiliation, metaphor, or older literary wording. This manual checking step was time-consuming, but it made the final gold labels more defensible.

5.4 Evaluation Setup

All five models received the same 100 inputs and the same binary instruction. ChatGPT and Gemini were called through OpenRouter because paid API access was available there. ALLaM, Jais, and Fanar were run locally through Ollama. The final run took 89.7 minutes, which made the experiment time- and effort-consuming, especially because local models were slow and had to be tested carefully.

Table 5.5: Models and execution method.

Report name	Exact model id	Execution method
ChatGPT	`openai/gpt-4o-mini`	OpenRouter API
Gemini	`google/gemini-2.5-flash`	OpenRouter API
ALLaM	`iKhalid/ALLaM:7b`	Ollama local run
Jais	`jwnder/jais-adaptive:7b`	Ollama local run
Fanar	`yasserrmd/Fanar-1-9B-Instruct`	Ollama local run

The final prompt was intentionally strict because the first prompt allowed too much freedom. The system prompt defined toxicity and required the model to output exactly one lowercase English word only: `yes` or `no`. The user prompt was:

```
Is the following Arabic text toxic?
[Arabic text]
Answer only yes or no:
```

For both API and local runs, the temperature was set to 0, the maximum output length was limited to 3 tokens, and the answer parser accepted only yes/no as valid. The final JSON log records one result per sample for each model.

5.4.1 Jais first-attempt failure and prompt correction

Jais caused a major problem in the first pilot run. With the earlier pseudo-JSON prompt, it frequently ignored the required output format and translated, paraphrased, repeated the sentence, or generated unrelated text. In that pilot output, Jais had 72/100 malformed or non-compliant outputs. This was not only a classification issue; it was a prompt-following failure.

After switching to the strict yes/no prompt, Jais produced valid yes/no outputs for all 100 final samples, so it is kept in the final ranking. However, its first-attempt failure is important because it shows that weaker local models may need much more constrained prompts.

Table 5.6: Examples of Jais non-compliance in the first pilot run. Text is redacted to avoid displaying offensive wording.

Pilot ID	Input type	Jais behavior in the pilot
1	Levantine toxic dialect sentence	Paraphrased the sentence instead of returning a label.
2	Gulf toxic dialect sentence	Repeated the input sentence instead of returning a label.
3	Gulf non-toxic sentence	Rewrote the sentence in MSA instead of returning a label.
4	Egyptian non-toxic greeting	Generated a greeting instead of returning a label.

5.5 Quantitative Results

Table 5.7 gives the final quantitative results over the full 100-case dataset. Gemini achieved the highest F1 score and accuracy. ChatGPT was second by F1, followed by ALLaM and Jais. Fanar was the most conservative model: it predicted non-toxic for many samples, which produced high precision but low recall.

Table 5.7: Overall final results for Arabic toxicity detection. Toxic is the positive class.

Model	Valid	Invalid	TP	TN	FP	FN	Precision	Recall	F1	Accuracy
ChatGPT	100	0	36	49	1	14	0.973	0.720	0.828	0.850
Gemini	100	0	40	50	0	10	1.000	0.800	0.889	0.900
ALLaM	100	0	37	45	5	13	0.881	0.740	0.804	0.820
Jais	100	0	40	38	12	10	0.769	0.800	0.784	0.780
Fanar	100	0	24	49	1	26	0.960	0.480	0.640	0.730

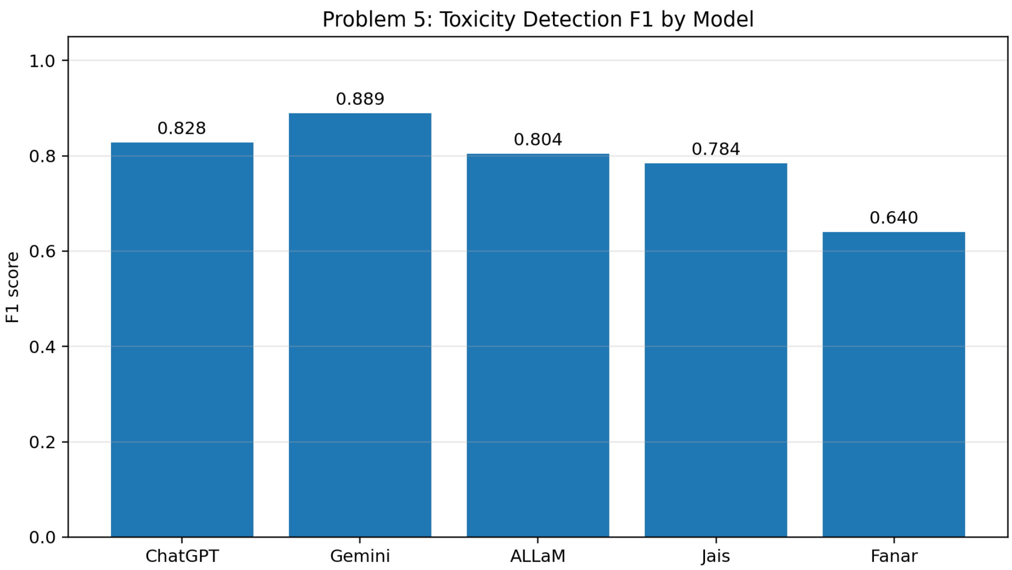

Figure 5.1: F1 score comparison across the five evaluated LLMs.

Table 5.8: F1 and accuracy by Arabic variety group.

Model	MSA F1	MSA Acc.	Classical F1	Classical Acc.	Dialect F1	Dialect Acc.
ChatGPT	0.875	0.882	0.815	0.848	0.786	0.818
Gemini	1.000	1.000	0.769	0.818	0.867	0.879
ALLaM	0.914	0.912	0.667	0.758	0.788	0.788
Jais	0.789	0.765	0.769	0.818	0.789	0.758
Fanar	0.759	0.794	0.316	0.606	0.741	0.788

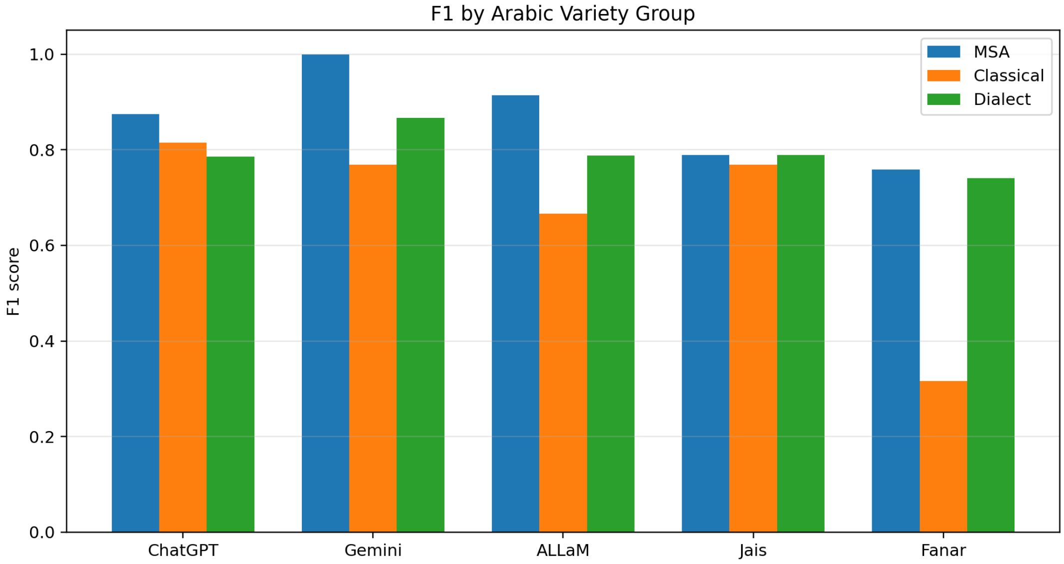

Figure 5.2: F1 score by Arabic variety group.

5.6 Selected Output Examples

Table 5.9 shows representative predictions from the final JSON log. To keep the printed report clean, the exact Arabic text is omitted here and can be found in the submitted CSV/JSON files. The examples include easy cases, false positives, false negatives, and cases where several models disagreed.

Table 5.9: Selected model outputs from the final run. Full Arabic inputs are in the submitted logs.

ID	Gold	ChatGPT	Gemini	ALLaM	Jais	Fanar	Wrong models
4	non-toxic	no	no	no	no	no	none
18	toxic	yes	yes	yes	yes	yes	none
5	non-toxic	yes	no	yes	no	yes	ChatGPT, ALLaM, Fanar
11	non-toxic	no	no	yes	yes	no	ALLaM, Jais
29	toxic	no	yes	no	yes	no	ChatGPT, ALLaM, Fanar
63	toxic	no	no	no	no	no	ChatGPT, Gemini, ALLaM, Jais, Fanar
85	toxic	no	no	no	no	no	ChatGPT, Gemini, ALLaM, Jais, Fanar
93	toxic	yes	yes	no	no	no	ALLaM, Jais, Fanar

The selected outputs show that obvious toxicity is generally easy for all models, while indirect hostility, dialect slang, and Classical Arabic satire are much harder. The biggest repeated issue is false negatives: several models recognize explicit modern insults but miss softer, indirect, or older forms of attack.

5.7 Qualitative Observations

Gemini had the best overall result with F1 = 0.889. It made no false positives, so it was very safe about not accusing non-toxic text. Its remaining errors were false negatives, especially in Classical Arabic and subtle dialect examples.

ChatGPT achieved F1 = 0.828. It had high precision (0.973) but lower recall (0.720), so it tended to miss toxic examples when the wording was indirect or literary.

ALLaM achieved F1 = 0.804. It was strong on MSA, where its F1 was 0.914, but it was weaker on Classical Arabic and produced some false positives on neutral dialect comments.

Jais achieved F1 = 0.784 in the final strict run. Its output format became valid, but it still produced many false positives. This means the strict prompt solved the first-attempt malfunction but did not fully solve the classification weakness.

Fanar achieved F1 = 0.640. Its precision was high (0.960), but its recall was only 0.480. It predicted no very often, causing many toxic samples to be missed.

5.8 Error Analysis

The assignment requires at least three incorrect or weak outputs, at least three Arabic-specific challenges, and at least two difficult or borderline cases. The following table gives representative errors from the final run. Offensive words are not printed; the IDs refer to the submitted logs.

Table 5.10: Representative incorrect or weak outputs. Text is summarized rather than quoted.

ID	Type	Gold	Model outputs	Explanation
5	MSA, conflict-related wording	non-toxic	ChatGPT=yes, Gemini=no, ALLaM=yes, Jais=no, Fanar=yes	False positive caused by negative-topic words. The sentence asks for less division, so the intent is non-toxic.
11	MSA, sympathy statement	non-toxic	ChatGPT=no, Gemini=no, ALLaM=yes, Jais=yes, Fanar=no	Neutral sympathy about victims; ALLaM and Jais treated sadness and death-related vocabulary as toxicity.
29	MSA, indirect metaphorical attack	toxic	ChatGPT=no, Gemini=yes, ALLaM=no, Jais=yes, Fanar=no	Some models missed toxicity when it appeared as metaphorical and political language rather than a direct modern insult.

ID	Type	Gold	Model outputs	Explanation
51	Gulf dialect, mixed affection and criticism	toxic	ChatGPT=no, Gemini=yes, ALLaM=yes, Jais=yes, Fanar=no	The sentence combines friendly wording with a redacted insult. ChatGPT and Fanar judged the overall tone as non-toxic.
63	Dialect/political sarcasm	toxic	ChatGPT=no, Gemini=no, ALLaM=no, Jais=no, Fanar=no	Subtle sarcastic attack; all five models missed it because the toxicity is pragmatic rather than explicit.
85	Classical Arabic satire/hija’	toxic	ChatGPT=no, Gemini=no, ALLaM=no, Jais=no, Fanar=no	Classical satire was treated as harmless literary text, although the intended function is hostile ridicule.
93	Classical humiliation	toxic	ChatGPT=yes, Gemini=yes, ALLaM=no, Jais=no, Fanar=no	The attack is expressed through old rhetorical style; local models especially missed the pragmatic humiliation.

5.8.1 Arabic-specific challenge 1: dialect variation

Dialectal Arabic contains local words, spelling variation, and social-media shortcuts that are not equally represented in model pretraining. The final dataset includes Gulf, Egyptian, Levantine, Iraqi, and Maghrebi items inside the dialect group. Several models missed subtle Gulf sarcasm, while some local models treated harmless dialect laughter or casual statements as toxic.

5.8.2 Arabic-specific challenge 2: morphology and orthographic variation

Arabic words can carry meaning through roots, patterns, affixes, and plural forms. Toxicity may appear through a derived form rather than a direct dictionary word. Informal spelling, missing hamza, final ta marbuta variation, and non-standard social-media orthography make exact keyword matching unreliable.

5.8.3 Arabic-specific challenge 3: ambiguity and context dependence

Some comments contain negative vocabulary but are not attacks, while others are toxic through sarcasm without using a clear slur. Sentence 5 discusses conflict but asks people to stop increasing division, so it is non-toxic. Sentence 63 is toxic mostly because of sarcastic framing and political attack. This shows that Arabic toxicity detection requires pragmatic understanding, not only word spotting.

5.8.4 Arabic-specific challenge 4: Classical Arabic satire

The Classical Arabic section was difficult because older insults and hija’ do not look like modern social-media abuse. Several models treated hostile literary lines as neutral because the style is poetic and archaic. Fanar was especially conservative on Classical

Arabic, with Classical F1 = 0.316. This suggests that modern Arabic safety classifiers may not capture older rhetorical forms of humiliation.

5.9 Difficult and Borderline Cases

Some items were difficult even for humans because they depend on social-media context, sarcasm, or literary background. Table 5.11 lists examples without printing the offensive text.

Table 5.11: Borderline or difficult cases.

ID	Type	Gold	Wrong	Why it is difficult
5	MSA negative-topic wording	non-toxic	3/5	It contains words related to conflict, but the communicative intent is to reduce division.
56	Context-dependent social-media item	toxic	4/5	The final label follows the reviewed benchmark label, but the offensive signal is weak without the original thread context.
63	Dialect/political sarcasm	toxic	5/5	The wording is sarcastic rather than a direct insult; models that depend on obvious slur words miss it.
85	Classical satire/hija'	toxic	5/5	Classical hostile poetry is hard because the offensive force is literary and archaic, not modern slang.

5.10 Limitations

The benchmark is useful but still small. It contains 100 examples, so a few borderline labels can noticeably change the metric. The dataset is balanced across the three required variety groups, but dialects inside the dialect group are not equally represented because the priority was to keep the total at 100 while using real examples.

A second limitation is context. Social-media toxicity often depends on previous comments, the target person, or a named group. Our inputs were single sentences or comments, so some cases lose context. This explains why a few original MPOLD labels look less obvious after cleaning and shortening.

A third limitation is model access. ChatGPT and Gemini required paid API access, while ALLaM, Jais, and Fanar were run locally through Ollama. Local inference was slow and resource-limited. The team had to wait for long local model runs, pay for API calls, modify the prompt, rerun the benchmark, and manually inspect labels and outputs. This made the work very time- and effort-consuming. We tried our best, became tired during repeated runs and manual checking, and still documented the failures honestly.

Finally, prompt format matters. Jais failed badly in the first attempt when the required answer format was less constrained. The final strict yes/no prompt improved compliance, but a different prompt or few-shot examples could change the classification results. Therefore, the results should be interpreted as a controlled zero-shot comparison under one strict prompt, not as a permanent ranking of the models.

5.11 Final Conclusions and Recommendations

Gemini was the best model in this benchmark, with F1 = 0.889 and accuracy = 0.900. ChatGPT came second, with very high precision but lower recall. Among local models, ALLaM was the strongest overall and performed especially well on MSA. Jais became format-compliant after the strict prompt but still made many false positives. Fanar was the most conservative model and missed many toxic examples, especially Classical Arabic attacks.

The main recommendation is to use stricter output prompts for local LLM evaluation. A simple yes/no format was much more robust than pseudo-JSON for Jais. For future work, the dataset should be expanded beyond 100 samples, dialects should be balanced inside the dialect group, and ambiguous MPOLD examples should be reviewed with more annotators. A few-shot prompt with examples of dialect sarcasm and Classical Arabic hija' may also improve recall. For production moderation, a model should not rely only on binary zero-shot LLM outputs; it should combine LLM classification with annotation guidelines, confidence thresholds, and human review for borderline cases.

5.12 Checklist

Table 5.12: Problem 5 checklist.

Requirement	Status	Evidence in our work
Task definition	Completed	Binary Arabic toxicity detection for Team 27 / Topic 27.
Dataset size	Completed	100 samples.
Balanced Arabic varieties	Completed	34 MSA, 33 Classical, 33 Dialect.
Real data majority	Completed	MPOLD social-media comments plus public-domain Classical Arabic.
Gold output for each sample	Completed	`gold_label` column in the final CSV.
At least two annotators	Completed	`annotator1` and `annotator2` included for every sample.
IAA reported	Completed	Simple agreement = 100.0%, Cohen's $\kappa = 1.000$.
Same prompt for all models	Completed	Strict yes/no prompt used for all five models.
All five required models	Completed	ChatGPT, Gemini, ALLaM, Jais, and Fanar evaluated.
Inputs, prompts, outputs logged	Completed	Final CSV and JSON evaluation logs are submitted with the report.
Metric reported	Completed	F1, precision, recall, accuracy, and confusion counts.
Comparison table	Completed	Table 5.7.
Visualization	Completed	F1 bar chart and variety-group F1 chart.
Selected outputs	Completed	Table 5.9.
Error analysis	Completed	Representative errors, Arabic-specific challenges, and borderline cases included.
Limitations and recommendations	Completed	Final sections discuss limitations and future improvements.

5.13 Submitted Files for Problem 5

The Problem 5 submission package should include the final dataset CSV, the final evaluation JSON log, the evaluation script/prompt used for the final run, this `prob5.tex` report section, and the generated figures. The important files are:

- `to_annotate.csv`: final 100-sample dataset with gold labels and annotator columns.
- `evaluation.json`: final outputs for ChatGPT, Gemini, ALLaM, Jais, and Fanar.
- `prob5.tex`: final report section.
- `Figures/problem5_f1_scores.png` and `Figures/problem5_f1_by_variety.png`: report visualizations.